



# Highly Efficient Neuromorphic Computing Systems with Emerging Nonvolatile Memories


Bonan Yan

Dept. Electrical & Computer Engineering  
Duke University

Slides available at: <https://bonanyan.github.io/bn/>


# Efficiency Is The Key to Ubiquitous AI

Limited Power/Energy



**How to build  
low-power &  
high-performance  
hardware?**

Better Accuracy Comes From Larger Models






- Memory is the Bottleneck; Data Movement Is Expensive
- How to overcome the memory bottleneck?

# Overhead Dominated by Memories

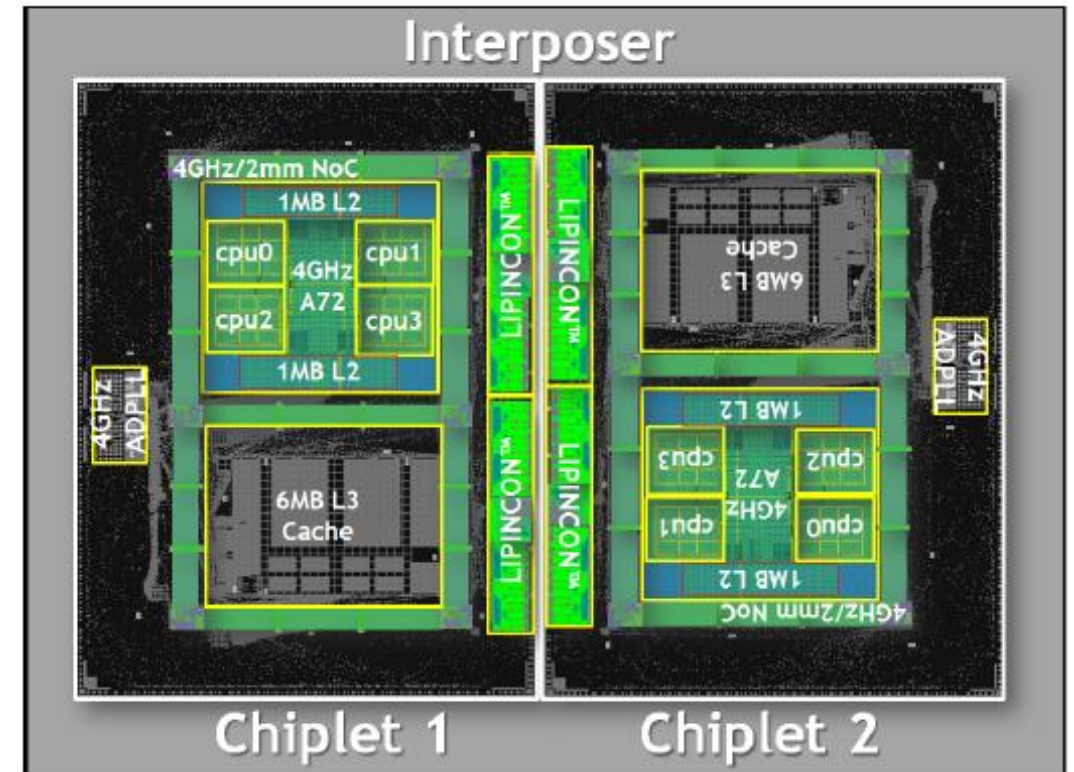
## Power



## Energy

Operations (32bit)	Energy (pJ)
Int ADD	0.1
Float ADD	0.9
Register File	1
Int Multiply	3.1
Float Multiply	3.7
SRAM Cache	5
DRAM Memory	640

## Area



# Specialized Hardware Enhances Efficiency

$$P = \alpha C V_{DD}^2 f$$





$P$ : Power Dissipation

$\alpha$ : Activity Factor


$C$ : Load Capacitance

$V_{DD}$ : Power Supply

$f$ : Clock Frequency



● MOSFET Scaling  
Transistor Size &  
Clock Frequency



● Multicore to  
Manycore

● Domain-Specific

groq™

Cambricon  
寒武纪


MYTHIC

habana  
An Intel Company

GRAPHCORE

NVIDIA

arm



Google TPU v3  
~200Watts  
(4 TPU to run AlphaGo,  
300x less power  
consumption)


Duke

Source: Vasileska et al., JTCN, 2008, Google, <https://www.youtube.com/watch?v=xL40UN5p6IY&t=673s>



# Uniqueness of Neural Network Execution

## Multiply-Accumulate (MAC):

streaming  $[x_1 \ x_2 \ \dots \ x_m]$   $\begin{bmatrix} w_{11} & \dots & w_{1n} \\ \vdots & & \vdots \\ w_{m1} & \dots & w_{mn} \end{bmatrix}$  stationary  $= [y_1 \ y_2 \ \dots \ y_n]$  outputs  $\Rightarrow y_j = \sum_{i=1}^m x_i \cdot w_{ji}$



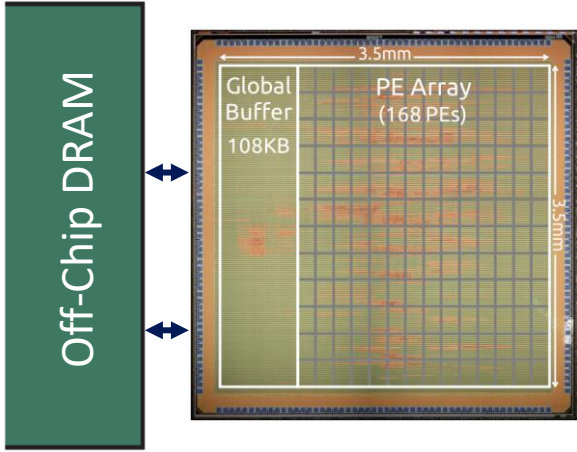
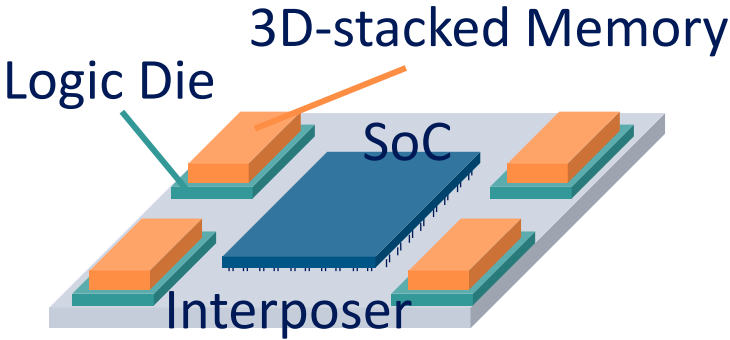
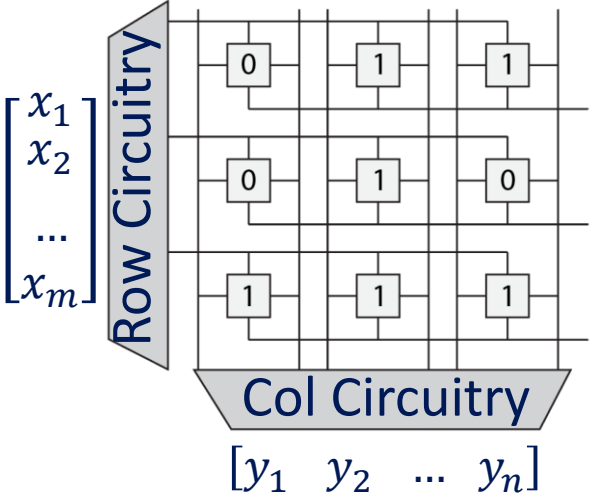
## Execution:



Inference:  
Inputs change;  
weights stay.  
How to fix weights  
to where MAC Units  
without moving?

- Need new circuitry to support in-memory computing
- Rethink memory from the ground up!

# Make Memory Access Less Expensive

	CMOS Accelerator (Planar Integration)	Near-Memory Computing	In-Memory Computing
Design Approach			
Energy to Access Weights		<p>200x~6x of MAC Energy</p> <p>55x~6x of MAC Energy</p>	<p>0</p>

Duke


DRAM

On-Chip Buffer

# Key Idea of In-Memory Computing

inputs weights outputs

$$\begin{bmatrix} V_1 & V_2 & V_3 \end{bmatrix} \begin{bmatrix} G_{11} & G_{12} & G_{13} \\ G_{21} & G_{22} & G_{23} \\ G_{31} & G_{32} & G_{33} \end{bmatrix} = \begin{bmatrix} I_1 & I_2 & I_3 \end{bmatrix}$$







- Weight Matrix Stored as Conductance  $G$
- Rely on Analog Computation (Kirchhoff's Current Law) for “almost free”
  - Multiplication:  $I = V \cdot G$
  - Addition:  $I^{column} = I_1^{row} + I_2^{row} + I_3^{row}$
- Ideal Nanoscale Devices for  $G$ :
  - Programmable Conductance
  - Multi-Level Cell
  - Small Footprint/High Density
  - Compatible with Existing CMOS Process



# Memristors for In-Memory Computing

Also Called **Resistive Random Access Memory**, RRAM or ReRAM

Cross-section TEM



Programmable resistor w/ analog states

ISSCC: Intel adds embedded ReRAM to 22nm portfolio

January 03, 2019

TSMC to start embedded RRAM production in 2019

According to reports, Taiwan Semiconductor Manufacturing Company (TSMC) is aiming to start producing embedded RRAM chips in 2019 using a 22 nm process. This will be initial "risk production" to gauge market reception.

	Multi-Level Cell	Cell Area	R/W Speed
SRAM	×	large	Fast
DRAM	×	medium	Medium
1T1R	✓	medium	Medium Fast
Flash	✓	small	Slow






# My work: Emerging Memory-Centric Design

- **Circuits & Systems Implementation**

- Spike-based Interface [DAC'15, DAC'18, DAC'20]
- Implementation of Neural Networks [VLSI'19, DAC'20]

- **Tolerate/Exploit Non-ideal Behavior of Memristors**

- Device Nonlinearity [ISCAS'16, IEDM'17, IEDM'19]
- Read Disturbance [ICCAD'17], Hard Fault [ITC'19]






# My work: Emerging Memory-Centric Design

- **Circuits & Systems Implementation**

- Spike-based Interface [DAC'15, DAC'18, DAC'20]
- Implementation of Neural Networks [VLSI'19, DAC'20]

- **Tolerate/Exploit Non-ideal Behavior of Memristors**

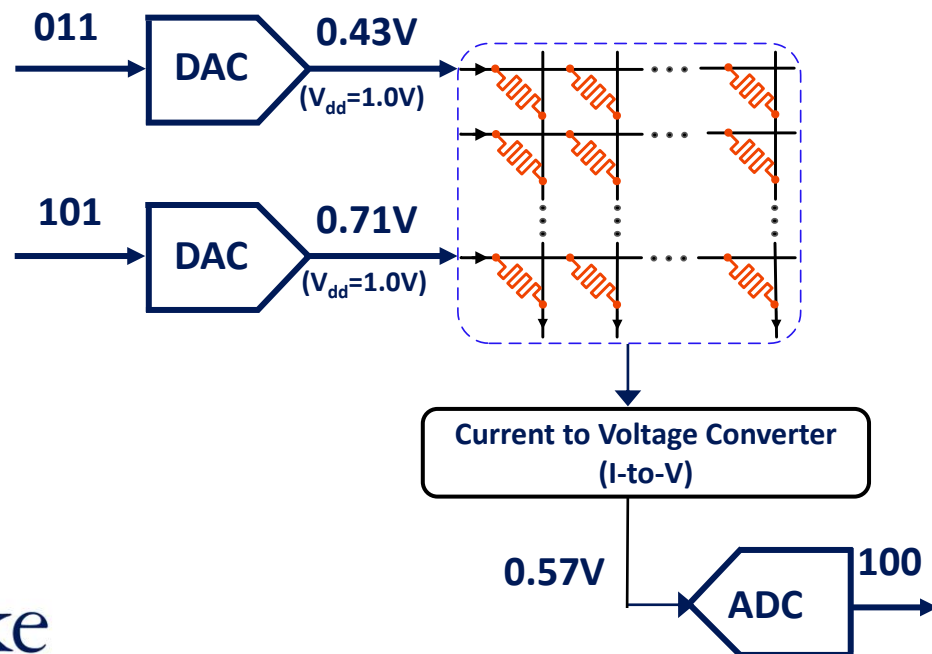
- Device Nonlinearity [ISCAS'16, IEDM'17, IEDM'19]
- Read Disturbance [ICCAD'17], Hard Fault [ITC'19]



# Conventional ADC is Too Large


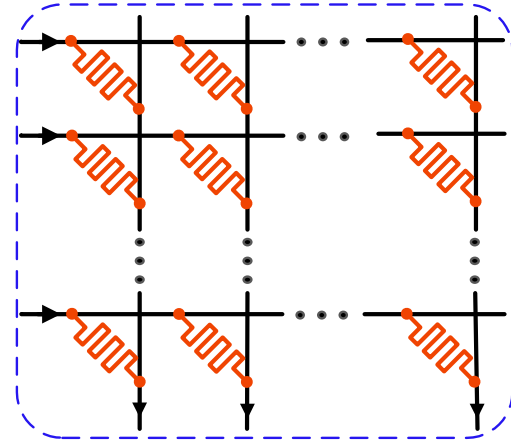
## The Level-based Design

- Compatible to existing signal processing
- High speed computation



Actual Layout:


256x256 1T1R Array



# My Approach: Spiking Interface Circuit

## What Better Designs Look Like


- Compute Parallely (Massive)
- Need Light-Weight Interface Circuitry



Duke


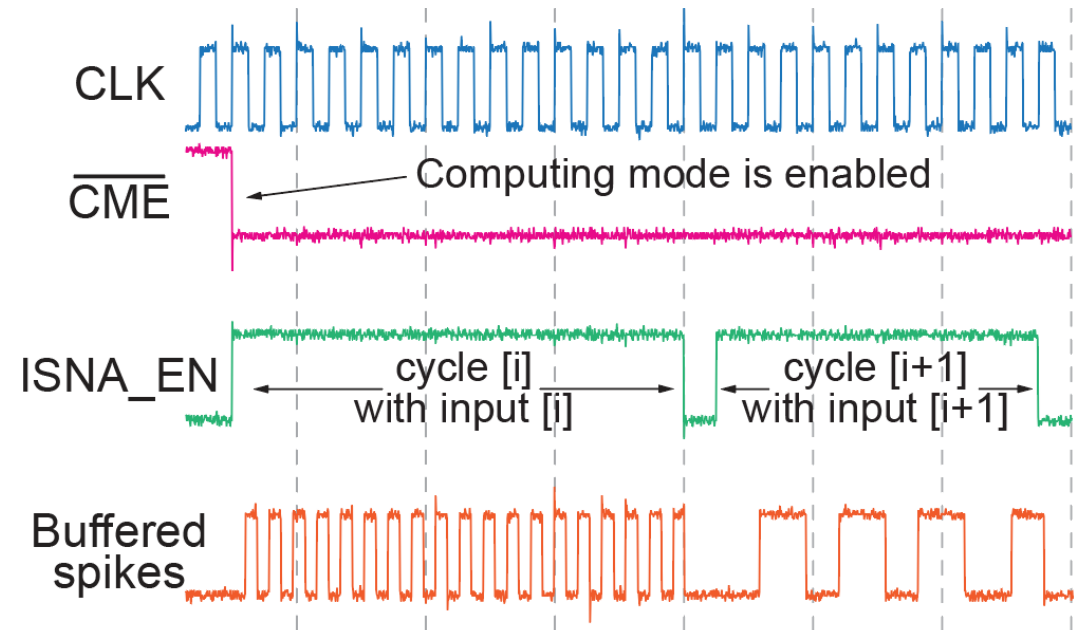
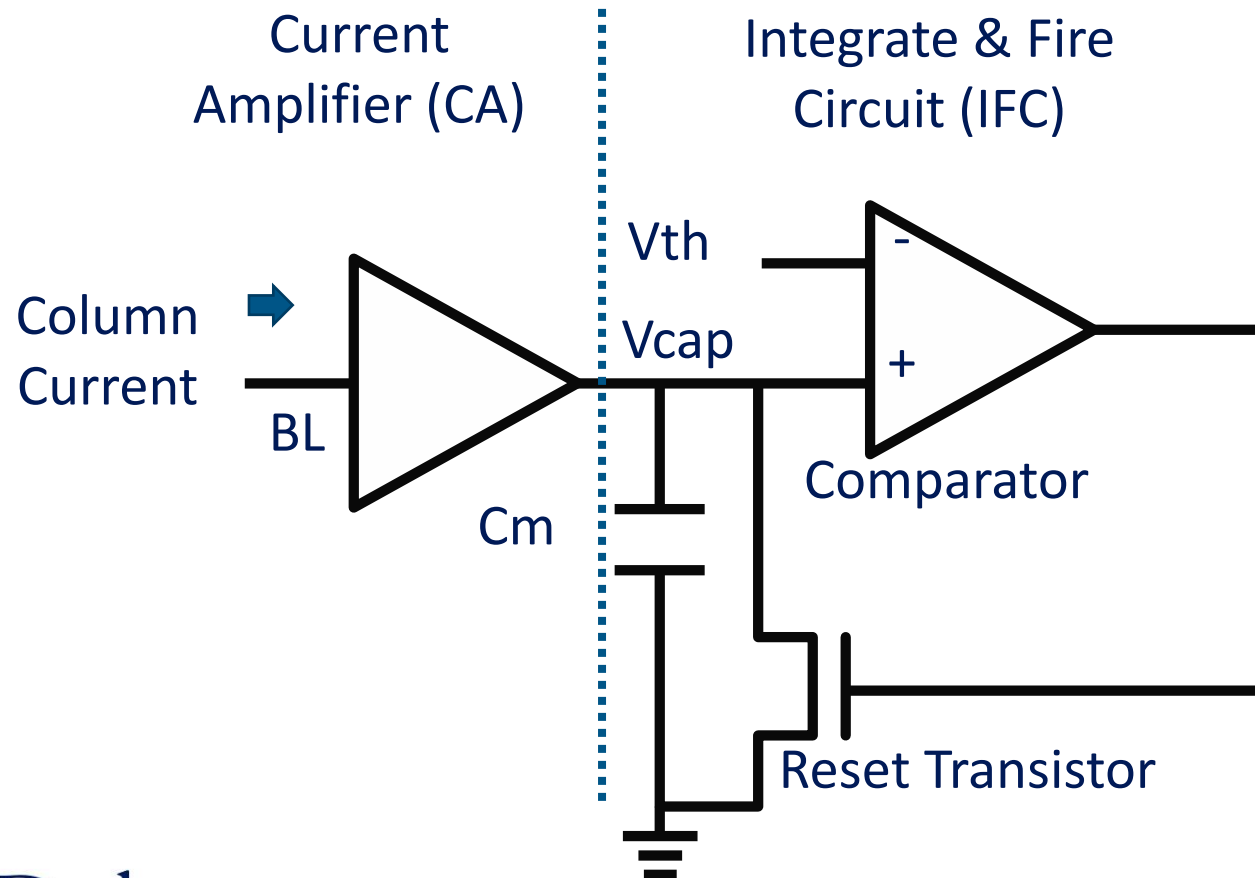
## The Spike-based Design

- Closer to biological system
- Extremely high power efficiency






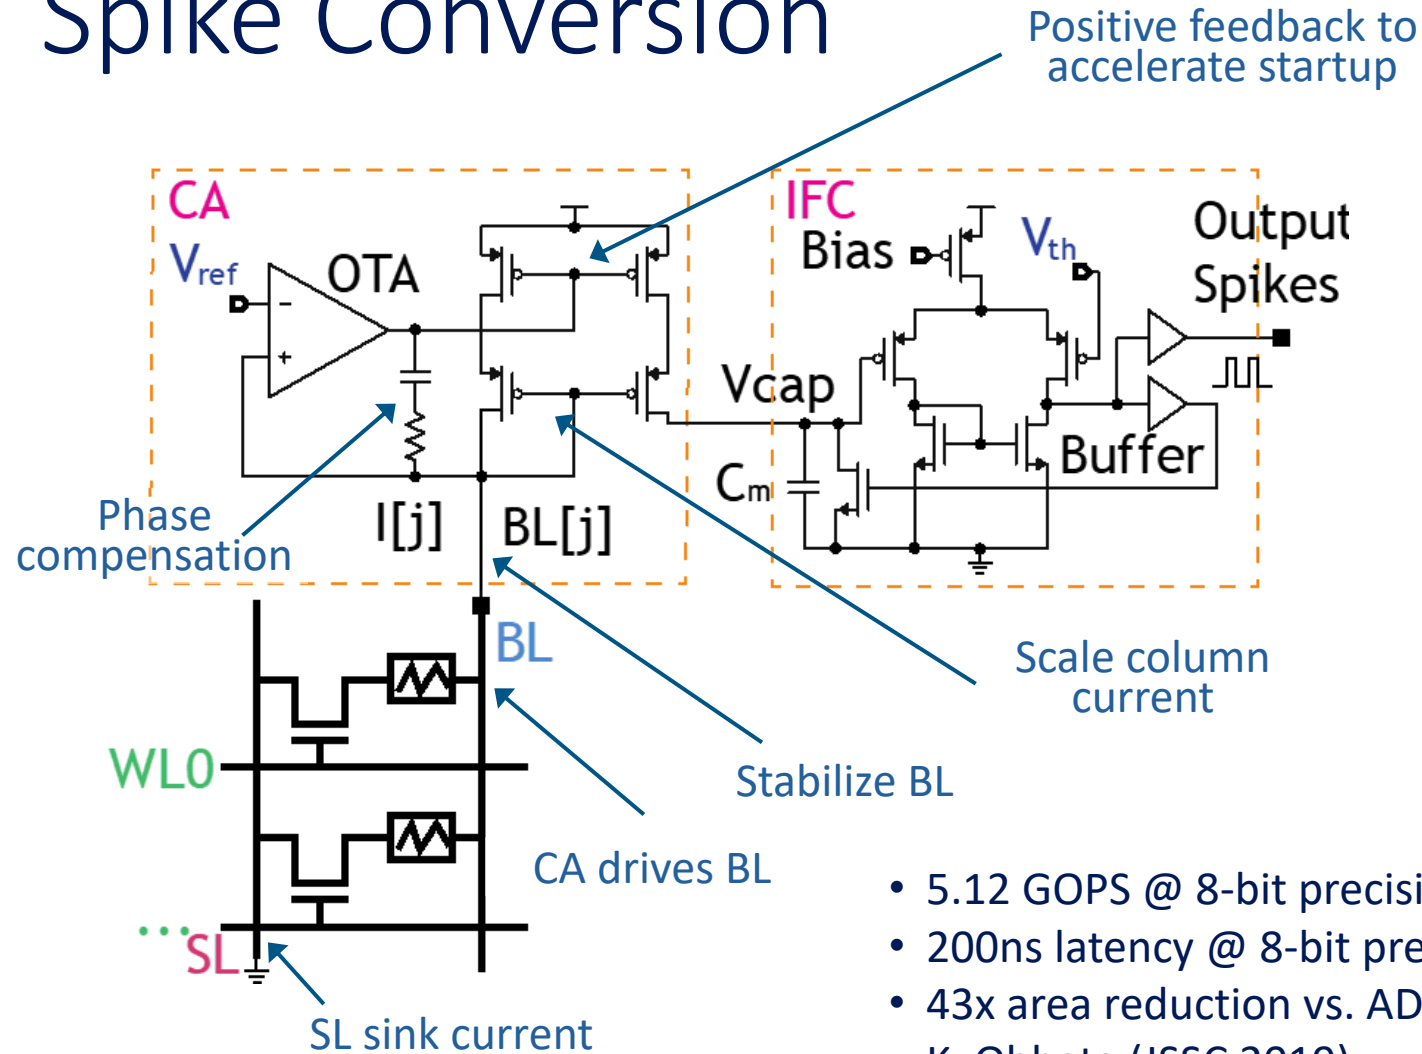
# Spike Conversion




Spike Conversion  
Circuit & Controller  
3152x3152  $\mu\text{m}^2$




# Spike Conversion



- 5.12 GOPS @ 8-bit precision
- 200ns latency @ 8-bit precision
- 43x area reduction vs. ADC by K. Ohhata (JSSC 2019)



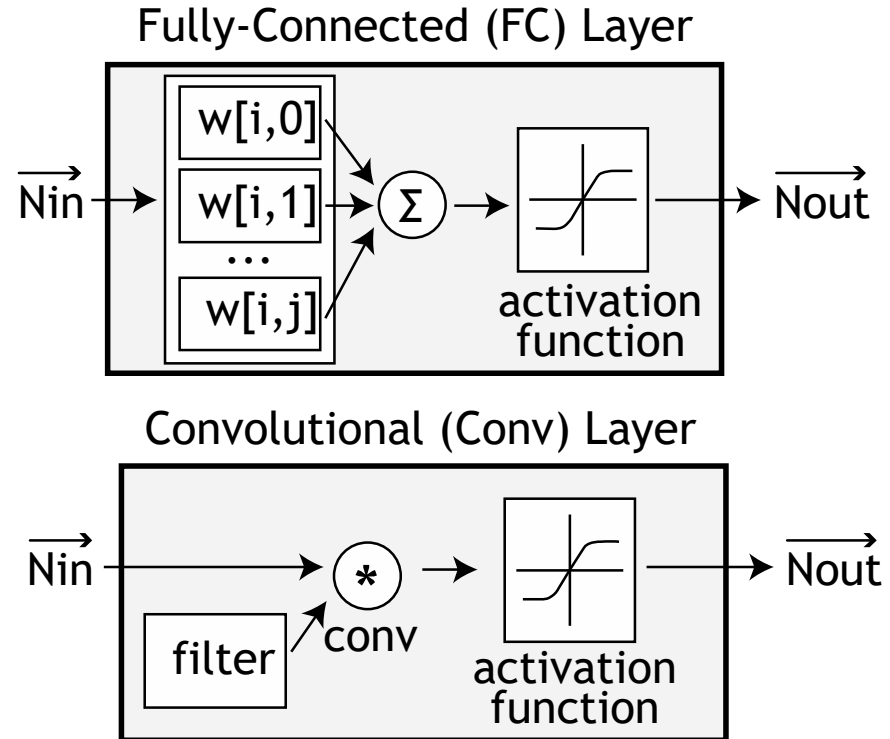
Tradeoff between large input current range and response speed



Enlarge phase margin tolerating capacitor positive feedback

# How to Use Spiking-Based Design to Execute Neural Networks?

# In Situ Nonlinear Activation (ISNA) Function



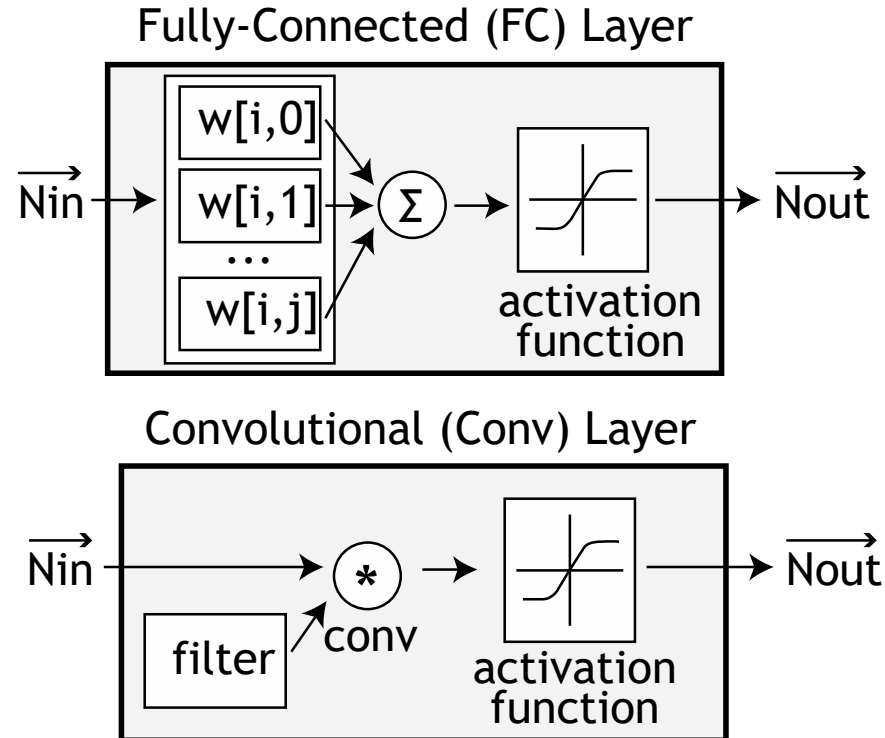
Single-layer Inference Operation:

- D** • Step 1: Load data from buffer
- A** • Step 2: Vector-matrix multiplication
- D** • Step 3: Nonlinear activation function
- D** • Step 4: Pooling
- D** • Step 5: Store results to buffer

**D** : digital domain      **A** : analog domain



# In Situ Nonlinear Activation (ISNA) Function



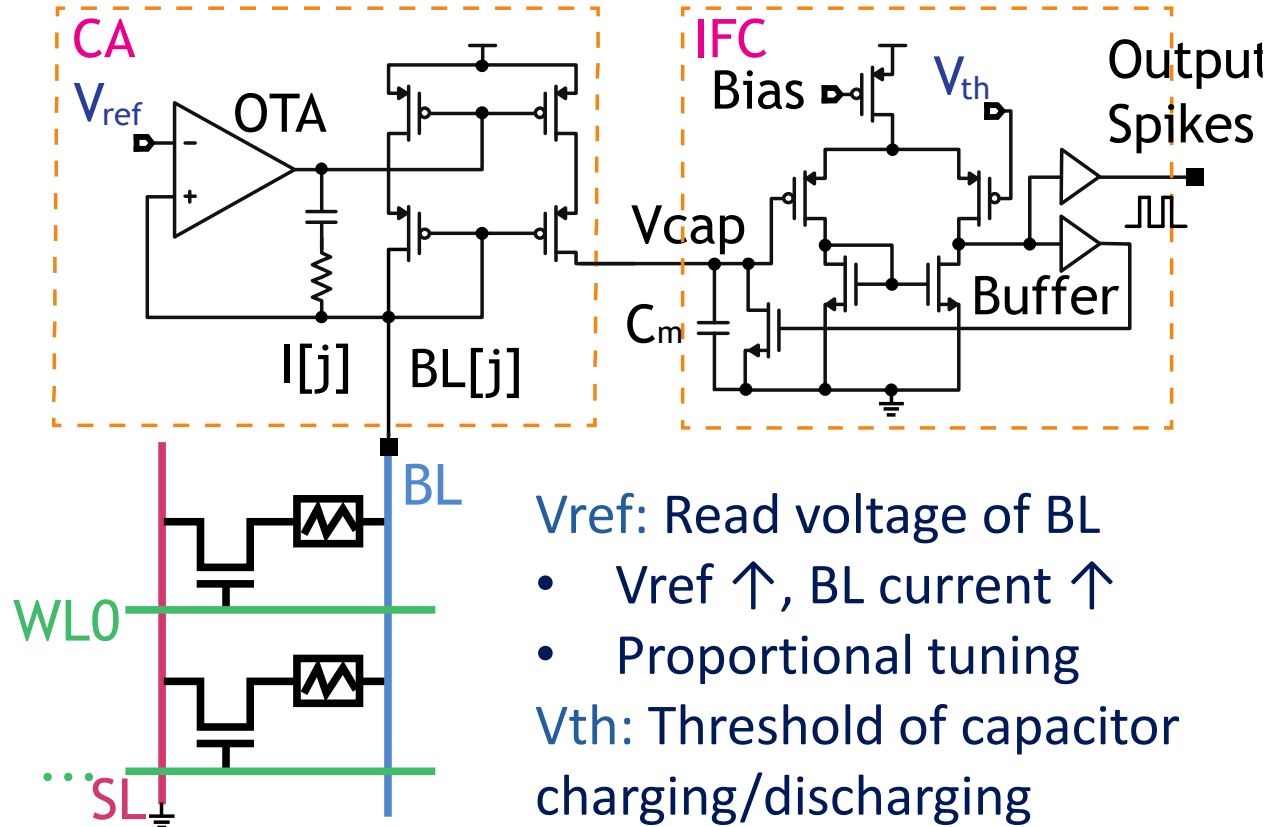
Single-layer Inference Operation:

- D** • Step 1: Load data from buffer
- A** • Step 2: Vector-matrix multiplication
- A** • Step 3: Nonlinear activation function
- D** • Step 4: Pooling
- D** • Step 5: Store results to buffer

**D** : digital domain      **A** : analog domain

Combine Step 2 & Step 3 to simplify PE operation:  
Use linear + nonlinear regions

# Adjust Activation Function

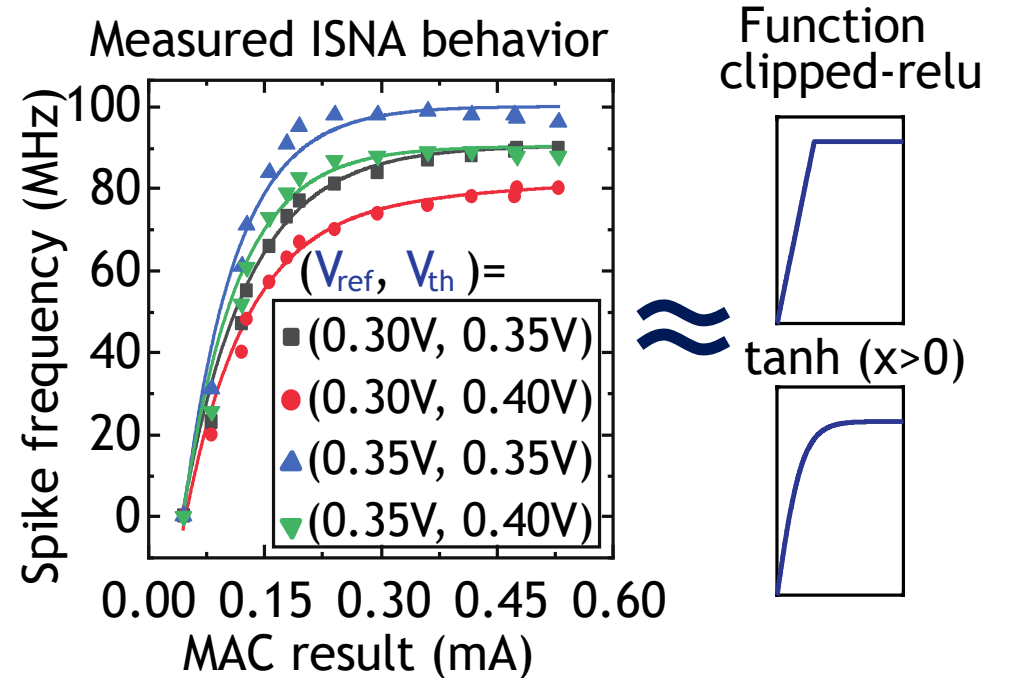


$V_{ref}$ : Read voltage of BL

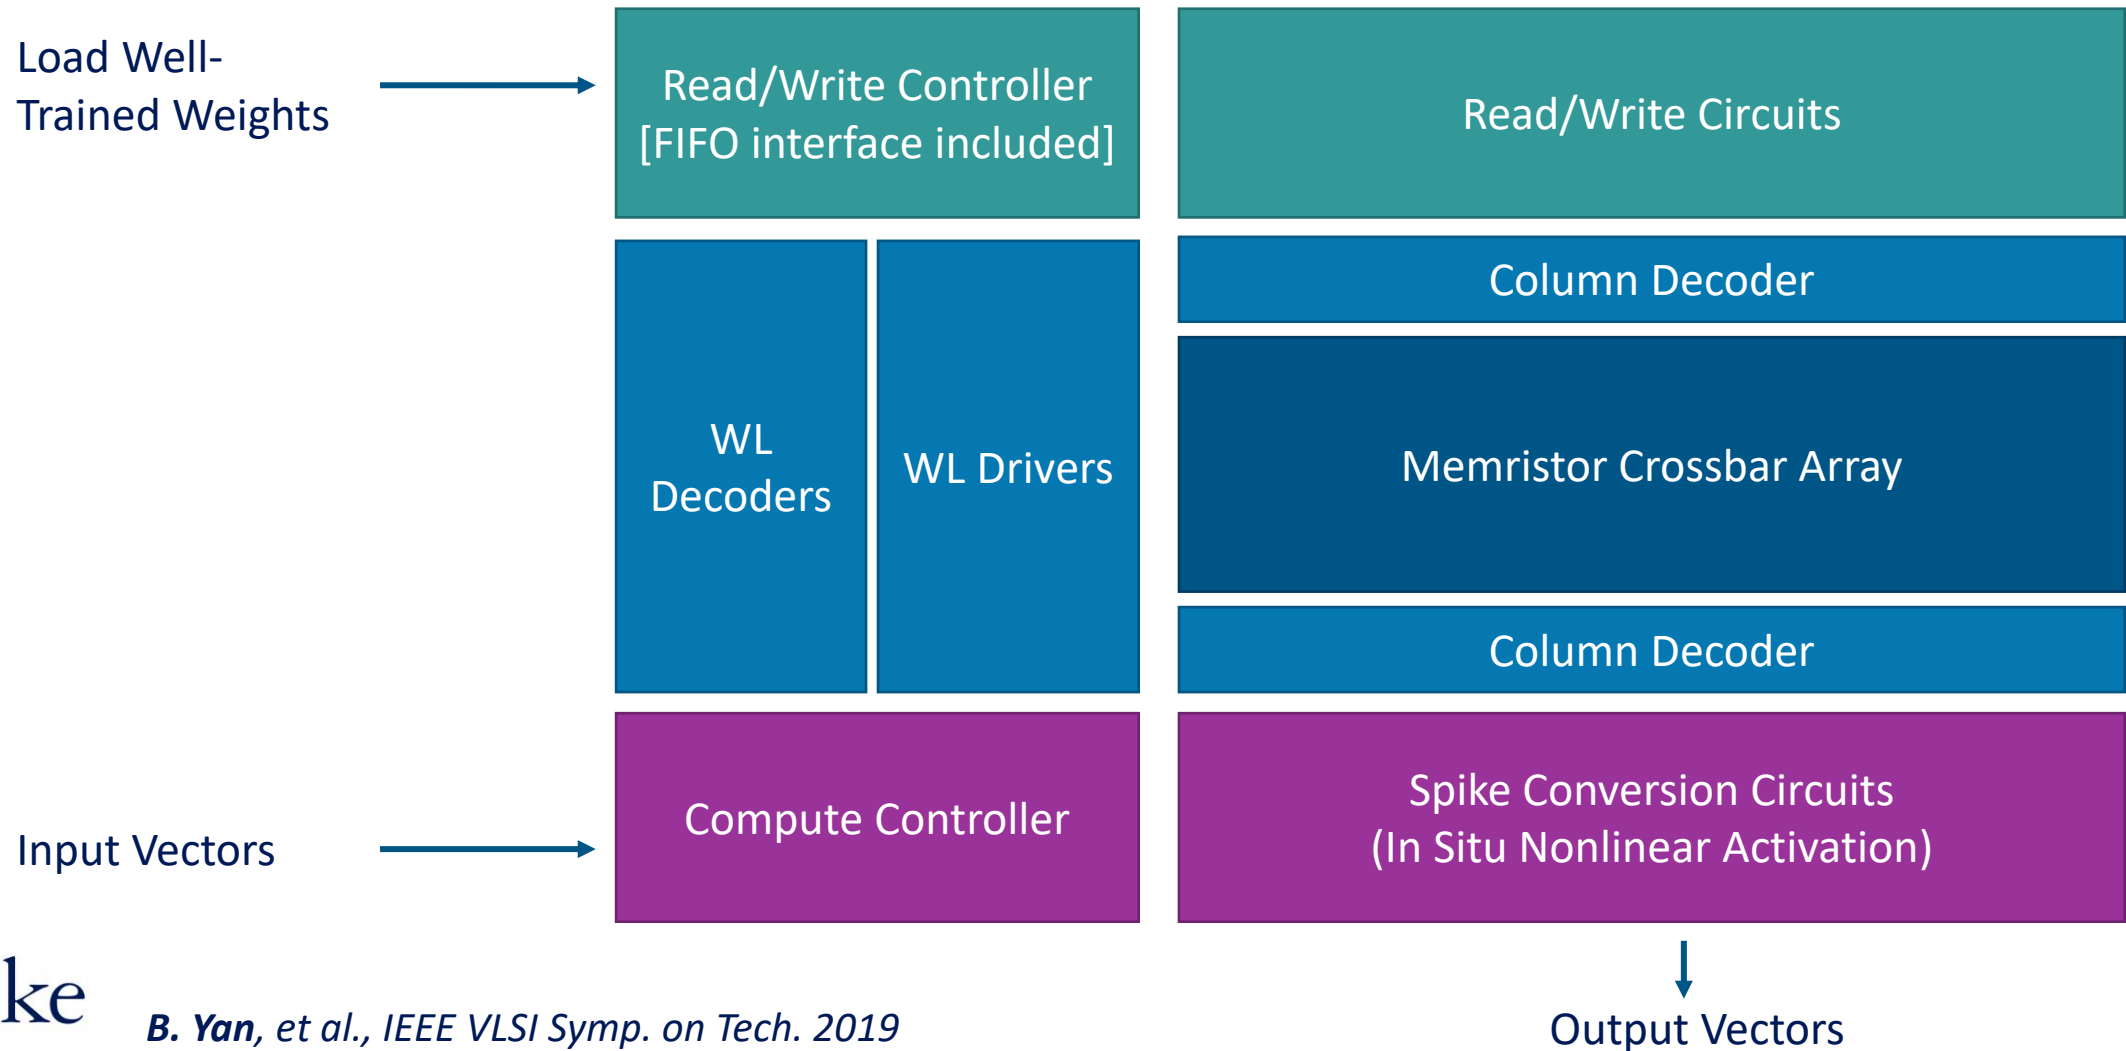
- $V_{ref} \uparrow$ , BL current  $\uparrow$
- Proportional tuning

$V_{th}$ : Threshold of capacitor charging/discharging

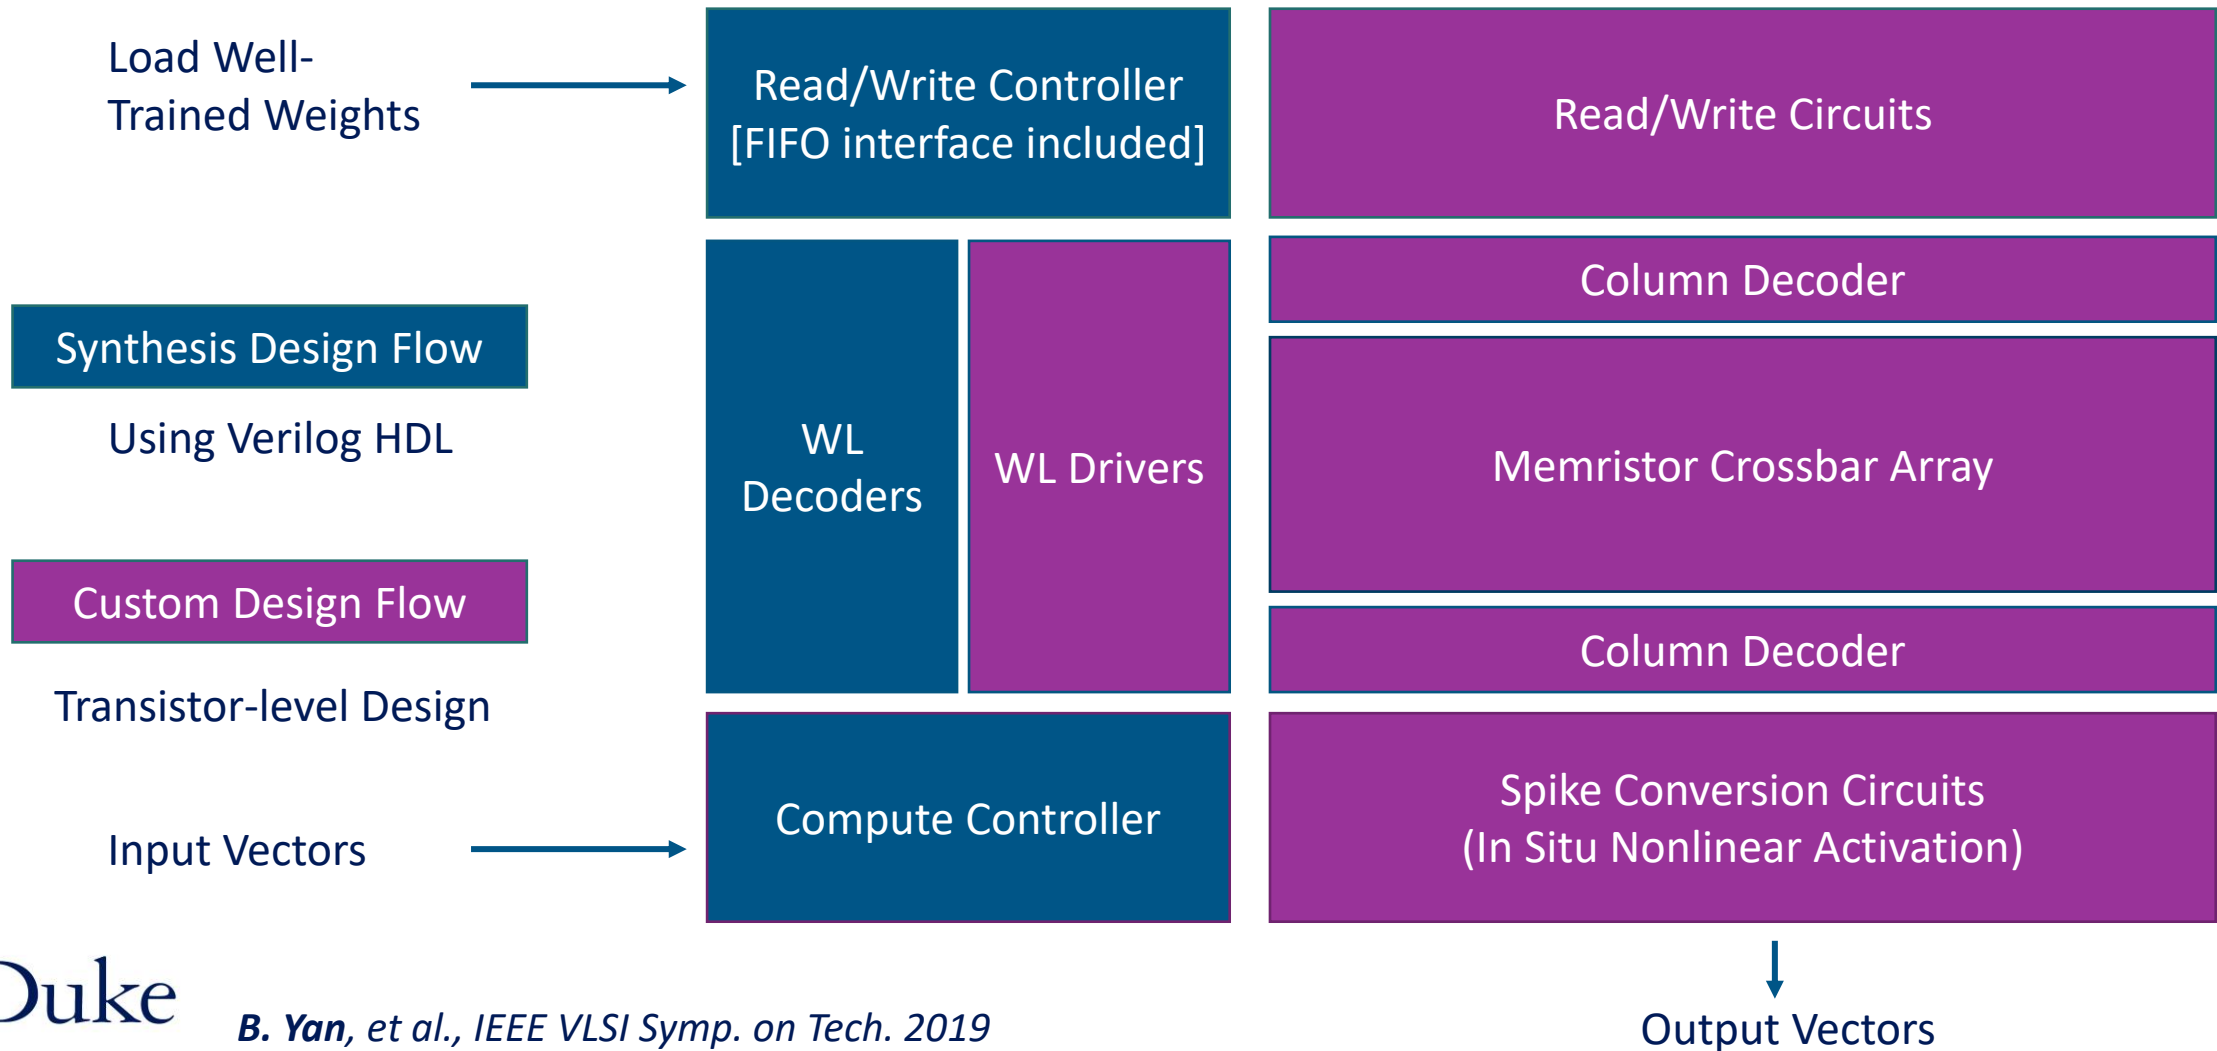
- $V_{th} \downarrow$ , Charging/discharging  $\uparrow$
- Distorted tuning



# Mixed-Signal Chip Architecture




# Mixed-Signal Design Flow






# Chip Summary



Technology	150nm CMOS +HfO <sub>x</sub> RRAM
Macro Capacity	64K (256×256)
Clock Frequency	50MHz
Energy Efficiency	0.257pJ/Mac
Average Power	1.52 mW
Layer-wise Latency	200ns
Real-time Benchmarks	3-layer perceptrons, LeNet-4, LetNet-5



GUI



Die Photo of RRAM Neuro-computing Engine



Camera

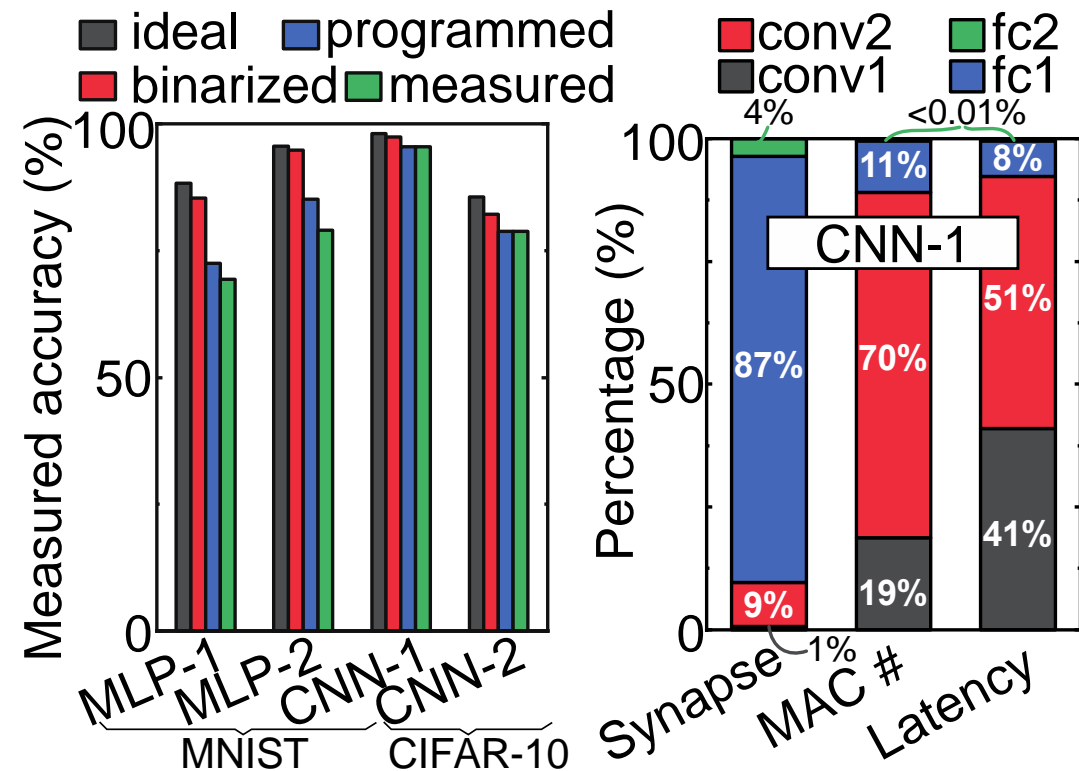
# Evaluation: Measured Neural Network Results

## MNIST:

- MLP-1: Single-layer perceptron
- MLP-2: 2-layer perceptron
- CNN-1: 4-Layer LeNet

## CIFAR-10:




- CNN-2: 5-Layer LeNet





# My work: Emerging Memory-Centric Design



- Circuits & Systems Implementation
  - Spike-based Interface [DAC'15, DAC'18, DAC'20]
  - Implementation of Neural Networks [VLSI'19, DAC'20]
- **Tolerate/Exploit Non-ideal Behavior of Memristors**
  - Device Nonlinearity [ISCAS'16, IEDM'17, IEDM'19]
  - Read Disturbance [ICCAD'17], Hard Faults [ITC'19]





# Non-Ideal Memristor - I

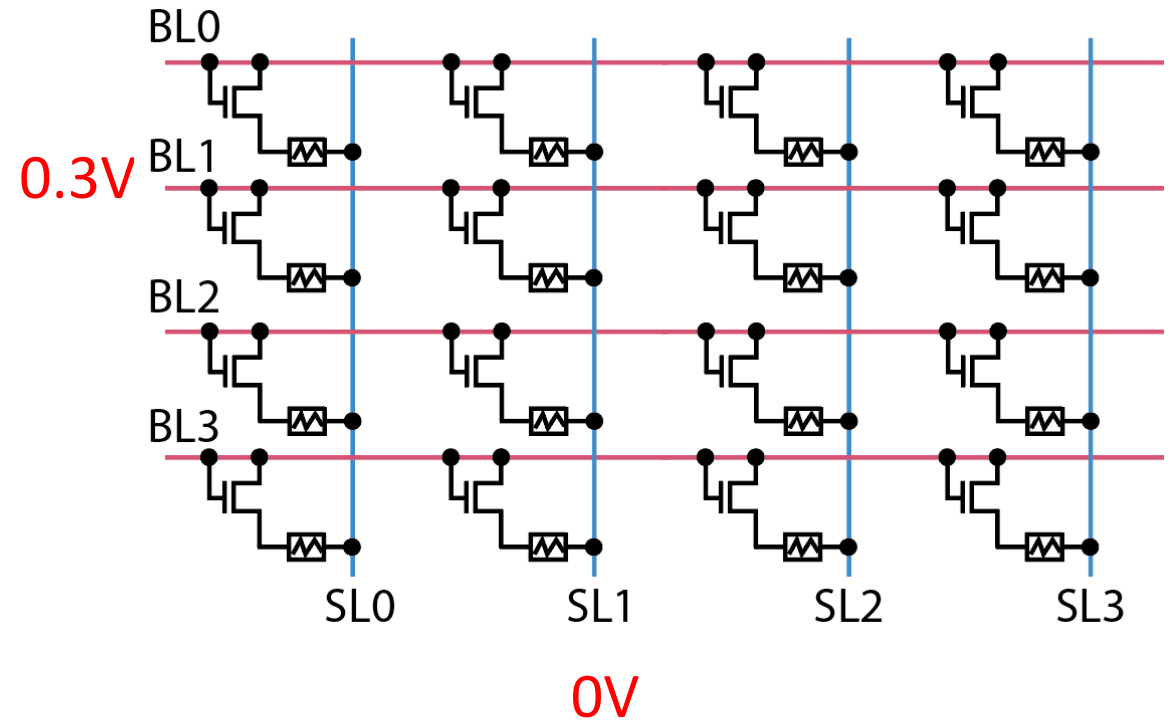

- Cell Nonlinearity: conductance varies when applied with different voltages



- Solution: Current Amplifier to Clamp Cell Voltage (shown in previous circuit design part)

# Non-Ideal Memristor - II


- Memristance Drift: conductance/memristance gradually deviates from original values under read voltage (read disturbance)
- Characteristic:
  - Very slow to observe



# How to Mitigate Memristance Drift for Inference?

$$E_{mse} = \sum_j^n \left( t_j - \sum_i^m w_{ij} x_i \right)^2 \quad E'_{mse} = \sum_j^n \left( t_j - \sum_i^m w_{ij} x_i - \sum_i^m \Delta w_{ij} x_i \right)^2$$

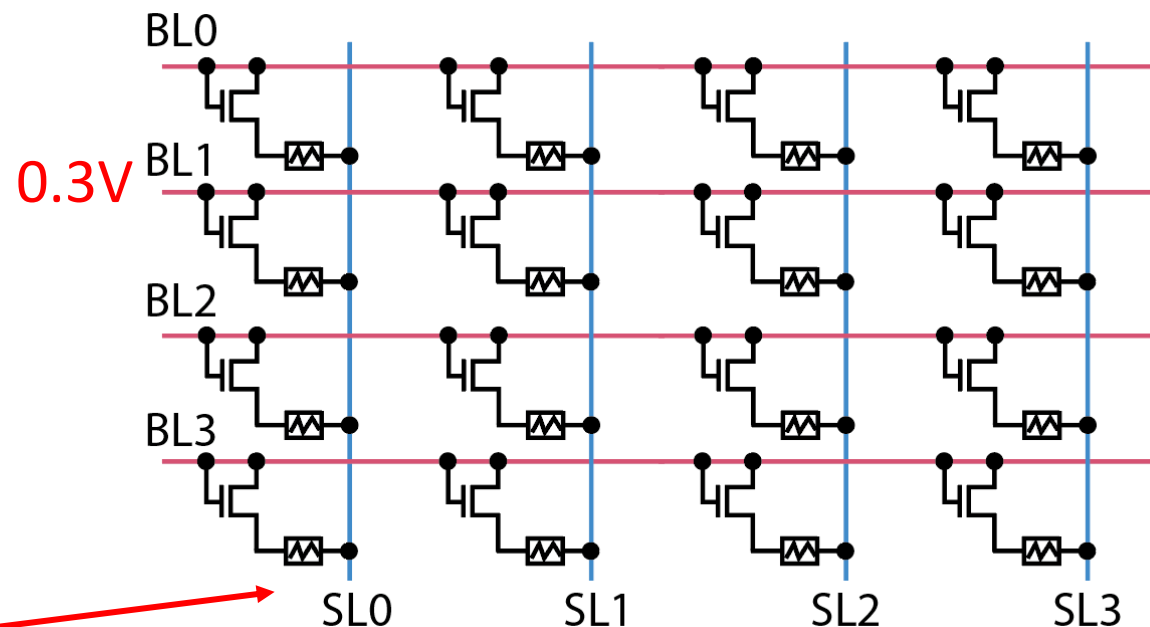
$E_{MSE}$ : MSE Error Function  
 $t_j$ : label for each neuron



$$\begin{aligned} \Delta E_{mse} &= E'_{mse} - E_{mse} \\ &= -2 \sum_j^n \left[ \left( t_j - \sum_i^m w_{ij} x_i \right) \left( \sum_i^m \Delta w_{ij} x_i \right) \right] \end{aligned}$$

Minimize Degraded Error Function:

$$\frac{\partial \Delta E_{mse}}{\partial \Delta w_{ij}} = -2 \underbrace{\left( t_j - \sum_i^m w_{ij} x_i \right)}_{\text{wavy red line}} x_i \leq 0$$

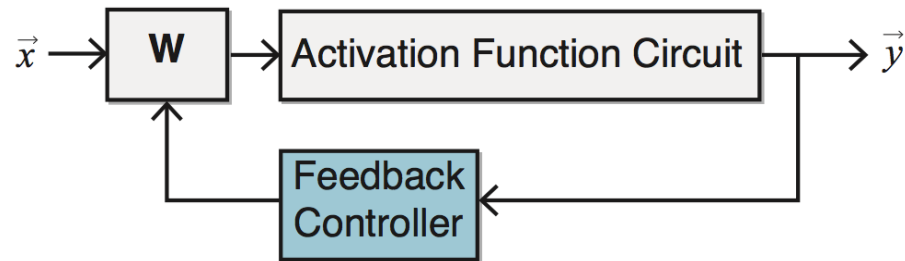


Change Column Current Direction


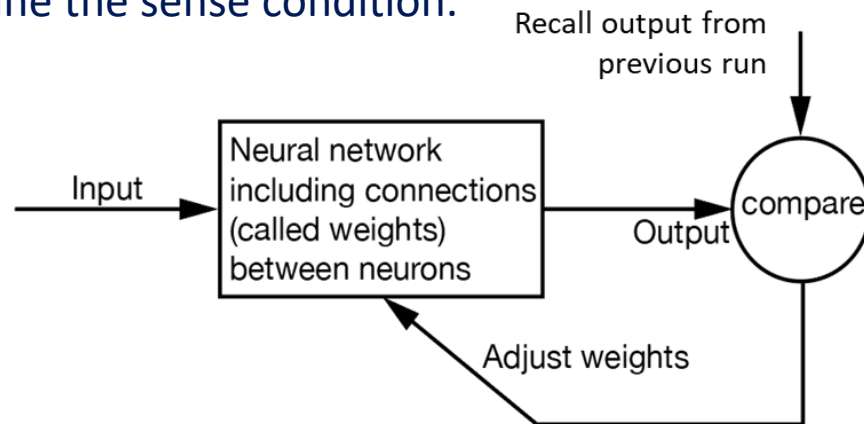
0V/0.6V

# Closed-loop Design to Enhance Weight Stability

**Feedback controller:** Adjust the voltage condition to compensate the memristance drift.




**“Arrogant principle”:** last output is used as the label to determine the sense condition.



# Non-ideal Memristor-III

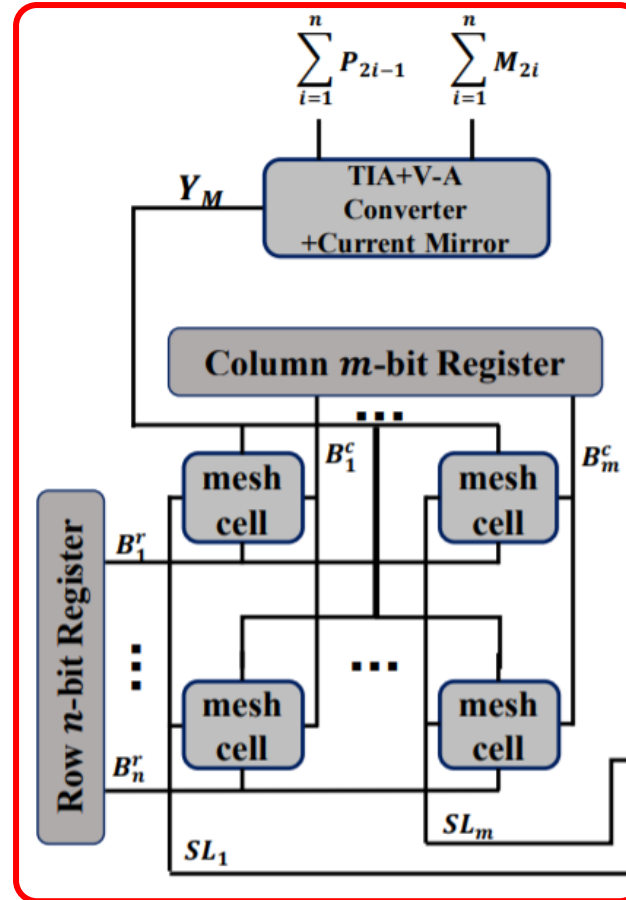
- Hard Faults: Dead Cells that are not programmable



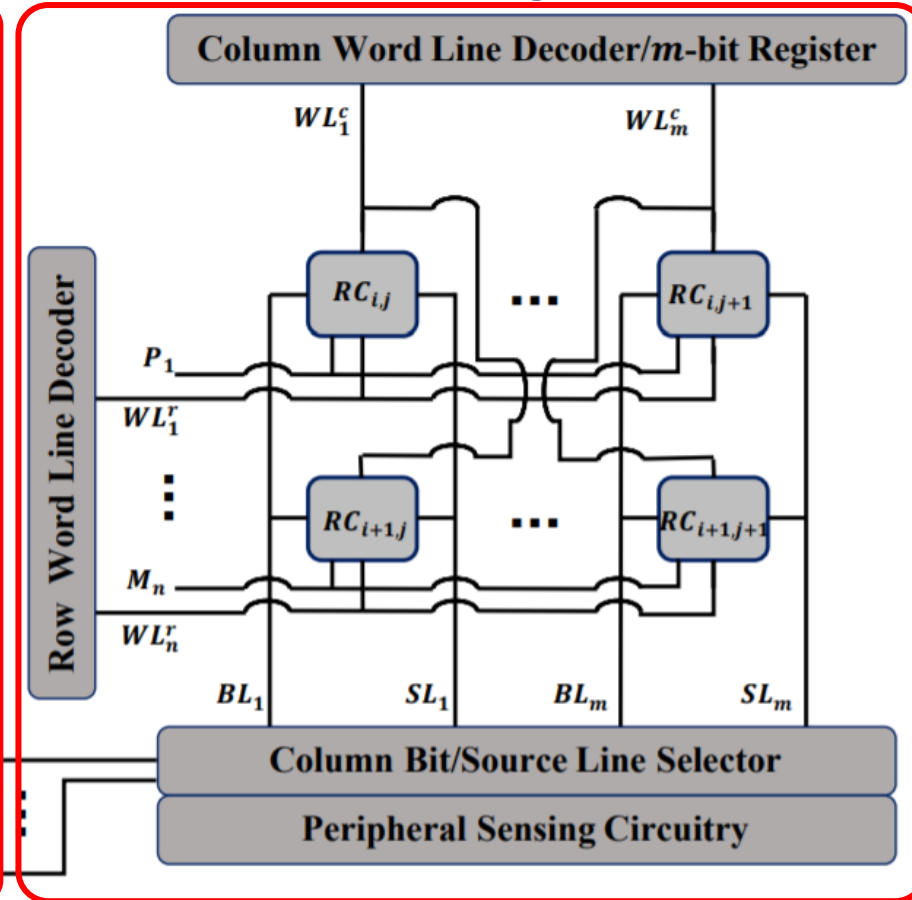
Noisy “Hewlett Packard”, 2016

- Sources of Hard Faults:
  - Fabrication
  - Beyond Endurance
- Our Solution:  
Error Correction Circuits

## Error Correction Circuit



## Core Engine



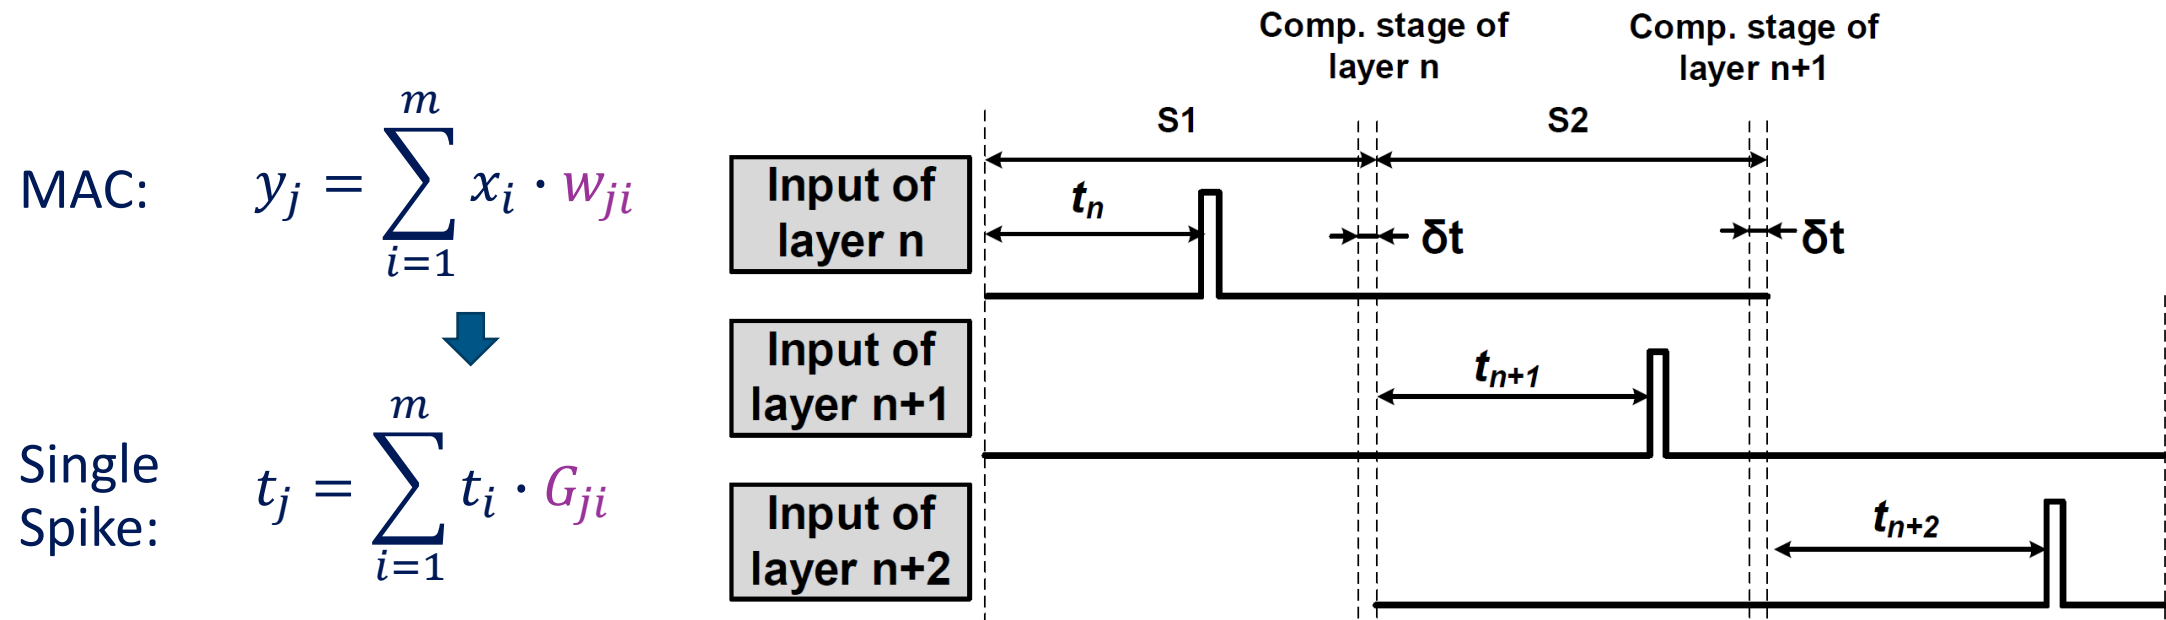
# Summary

- The Past and Now of AI Hardware
  - In-Memory Computing Eliminates Weights Movement
- My Work:
  - Spiking-based In-Memory Computing Engine Offer Very High Energy Efficiency & Good Performance
  - Clever Design Methodologies (e.g., Closed-Loop Sensing) is Effective to Tolerate Nonideal Features of Memristors



# Future Work I: Single Spike Processing Engine

Use Single Spike to Replace Multiple Spikes: ~10x Improvement of Energy Efficiency



“ReSiPE: ReRAM-based Single-Spiking Processing-In-Memory Engine”

Simulation Results Coming in July at Design Automation Conference (DAC) 2020


Duke

Expect to IC Tape-Out + Single-Spike Encoding Method [Potential Collaboration Opportunity]

# Future Work II: In-Memory Compute & Cache

## Challenge:

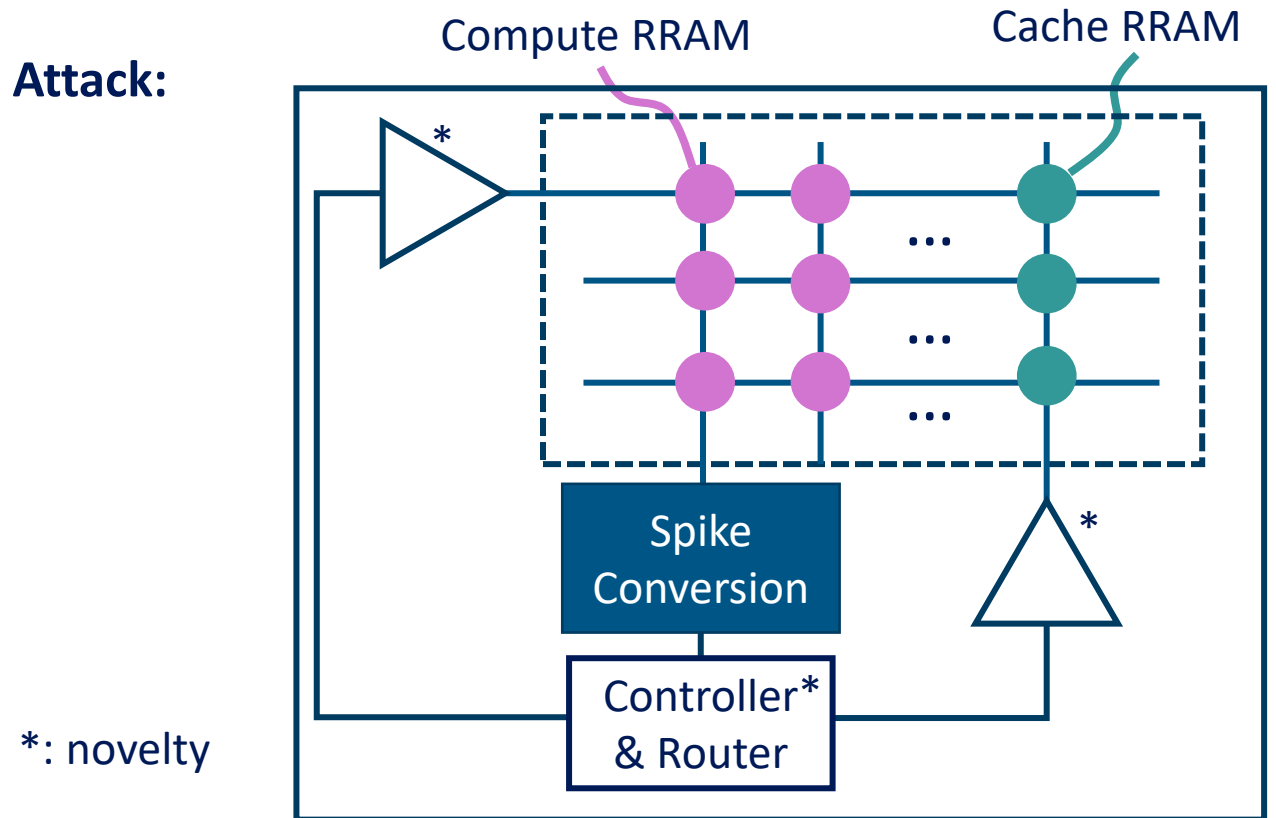
During Training:



No Need to Transport Weight

External DDRAM Access (Expensive!)

## Attack:



\*: novelty

## Prospect Applications:


NN Training Acceleration

Self-Updatable In-Memory Computing Engine


# Long-Term Prospects

## Compute


Rethink Compute  
Memory Hierarchies



RRAM




PCM



MRAM


“Organ” for Compute &  
Cache w/ different  
emerging memory types

In Situ Data Processing  
(Near-Sensor Computing)



## Control

Analog Content  
Addressable Memories



Fuzzy Logic Hardware

# Many Thanks for the Support!

## Advisors:

Prof. **Hai “Helen” Li**

Prof. **Yiran Chen**



## Collaborators:

Prof. **Joe Qiu**

Army Research Office

Dr. **Qing Wu**

Air Force Research Laboratory

Prof. **Jianhua (Joshua) Yang**

UMASS Amherst

Prof. **Meng-Fan Chang**

National Tsing-Hua Univ.



Prof. **Krishnendu Chakrabarty**

Duke University

Prof. **Weisheng Zhao**

Beihang University

## Student Collaborators:

**Ziru Li, Qilin Zheng, Brady Taylor**

Duke



Qualcomm



Thanks for Listening & Qs!

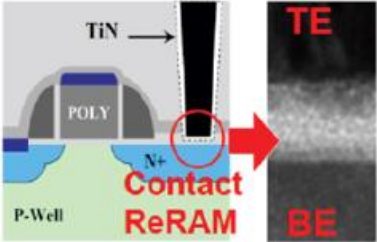
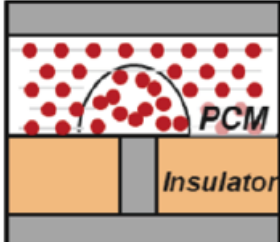
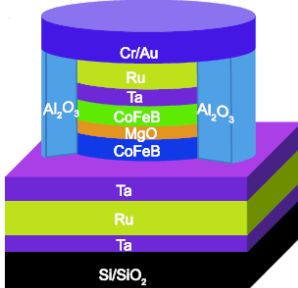
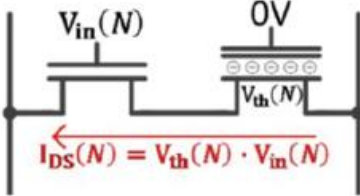
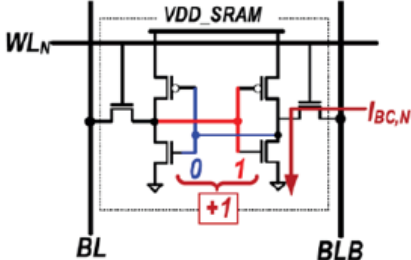
slides available at:

<https://bonanyan.github.io/bn/>

# Backup Slides



# In-Memory Computing Candidates

RRAM/Memristor	Phase-Change Memory	MRAM	Flash	SRAM
				
Nonvolatile	Nonvolatile	Nonvolatile	Nonvolatile	Volatile
Dense Cell Area	Dense Cell Area	Dense Cell Area	Dense Cell Area	Large Area
Bipolar Write	Unipolar Write	Bipolar Write	Erase-Needed Write	Bipolar Write
Fast Write	Fast Write	Fast Write	Slow Write	Ultra-Fast Write
Multi-level Cell	Multi-level Cell	Single-level Cell	Most Single-level Cell	Single-level Cell

# Comparison

Type		CIM Macro				CIM PE		Digital Processor
Work		VLSI'18 Panasonic	ISSCC'18 NTHU	ISSCC'18 NTHU	ISSCC'18 MIT	This work <sup>†</sup>	ISSCC'18 UIUC	TrueNorth [10]
Technology		180nm	65nm	65nm	65nm	150nm	65nm	28nm
Synapse		1T1R RRAM	1T1R RRAM	6T-SRAM	10T-SRAM	1T1R RRAM	6T-SRAM	SRAM
Nonvolatility		Yes	Yes	No	No	Yes	No	No
Standby current		~zero	~zero	high	high	~zero	high	high
Spiking NN		No	No	No	No	Yes	No	Yes
Capacity		2M	1M	4K	16K	64K	128K	256M
Cell area [F <sup>2</sup> ]		—	59	124	968	74	~256	—
Normalized die area		12×	—	—	30×	1×	11×	~17240×
Chip power [mW]		15.8	—	—	—	1.52	—	204.4
Activation precision		1 bit	3 bit	1 bit	7 bit	1~8 bit	8 bit	1 bit
Power efficiency [TOPS/W]	MAC only	20.7	16.95	55.8	28.1	—	—	—
	MAC +Activation	—	—	—	—	16.9	3.125	0.4
On-chip Activation Function Integration		No	No	No	No	Yes (tanh)	Yes (relu)	Yes (relu)
FoM*		—	0.86	0.45	0.20	1.83	0.098	—