



Highly Efficient Neuromorphic Computing Systems with Emerging Nonvolatile Memories


Bonan Yan

Dept. Electrical & Computer Engineering
Duke University

Slides available at: <https://bonanyan.github.io/bn/>


Efficiency Is The Key to Ubiquitous AI

Limited Power/Energy




**How to build
low-power &
high-performance
hardware?**

Better Accuracy Comes From Larger Models




Overhead Dominated by Memories

Power




Area



- **Memory is the Bottleneck; Data Movement Is Expensive**
- **How to overcome the memory bottleneck?**

Energy

MEM ADD	0.1
Float ADD	0.9
Register File	1
Int Multiply	3.1
Float Multiply	3.7
SRAM Cache	5
DRAM Memory	640



- : Function Unit
- : On-chip Memory
- : Control Module

Specialized Hardware Enhances Efficiency

$$P = \alpha C V_{DD}^2 f$$





P : Power Dissipation

α : Activity Factor


C : Load Capacitance

V_{DD} : Power Supply

f : Clock Frequency



● MOSFET Scaling
Transistor Size &
Clock Frequency



● Multicore to
Manycore

● Domain-Specific

groq™

MYTHIC


Cambricon
寒武纪

habana
An Intel Company

GRAPHCORE

NVIDIA

arm



Google TPU v3
~200Watts
(4 TPU to run AlphaGo,
300x less power
consumption)


Duke

Source: Vasileska et al., JTCN, 2008, Google, <https://www.youtube.com/watch?v=xL40UN5p6IY&t=673s>



Uniqueness of Neural Network Execution

Multiply-Accumulate (MAC):

streaming $[x_1 \ x_2 \ \dots \ x_m]$ $\begin{bmatrix} w_{11} & \dots & w_{1n} \\ \vdots & & \vdots \\ w_{m1} & \dots & w_{mn} \end{bmatrix}$ stationary $= [y_1 \ y_2 \ \dots \ y_n]$ outputs $\Rightarrow y_j = \sum_{i=1}^m x_i \cdot w_{ji}$



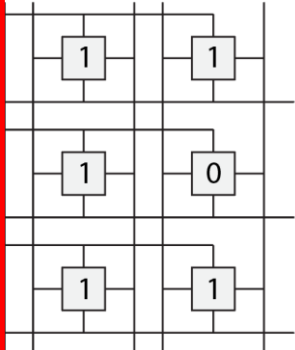


Execution:



Inference:
Inputs change;
weights stay.
How to fix weights
to where MAC Units
without moving?

Make Memory Access Less Expensive

	CMOS Accelerator (Planar Integration)	Near-Memory Computing	In-Memory Computing
Design Approach	<div>   </div> <p>High Bandwidth Memory</p>		 <p>Col Circuitry [y_1 y_2 ... y_n]</p>
Energy to Access Weights	<p>200x~6x of MAC Energy</p> <p>55x~6x of MAC Energy</p>		0

Duke

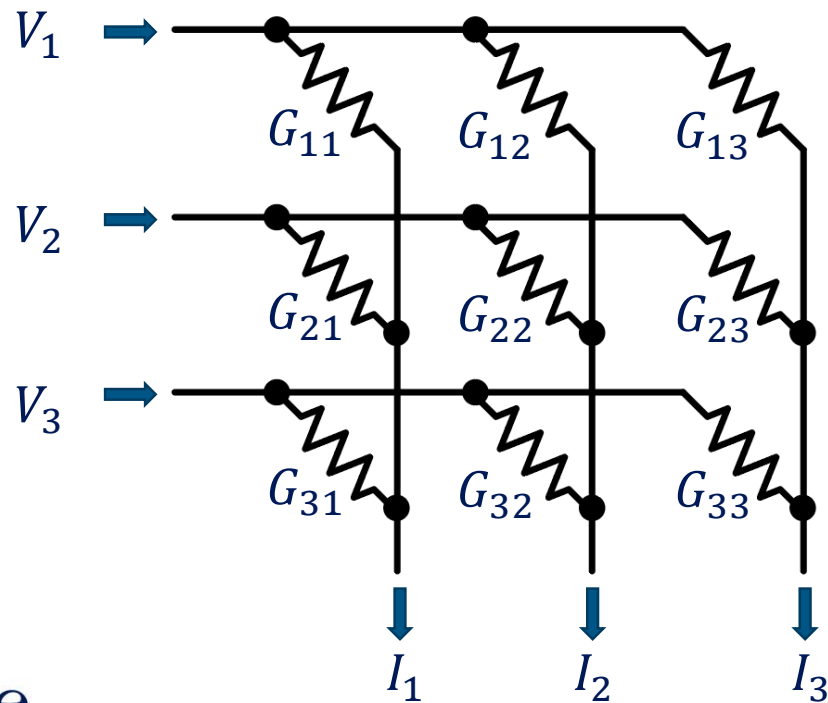
DRAM

On-Chip Buffer

Key Idea of In-Memory Computing

inputs weights outputs

$$\begin{bmatrix} V_1 & V_2 & V_3 \end{bmatrix} \begin{bmatrix} G_{11} & G_{12} & G_{13} \\ G_{21} & G_{22} & G_{23} \\ G_{31} & G_{32} & G_{33} \end{bmatrix} = \begin{bmatrix} I_1 & I_2 & I_3 \end{bmatrix}$$

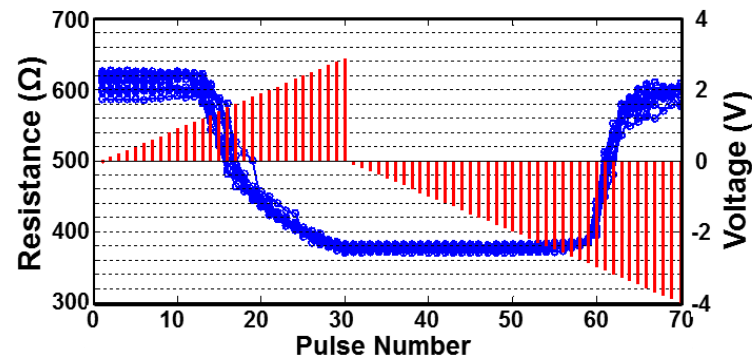
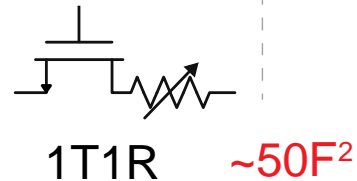
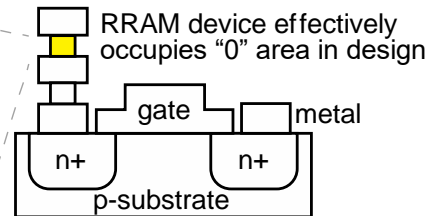



- Weight Matrix Stored as Conductance G
- Rely on Analog Computation (Kirchhoff's Current Law) for “almost free”
 - Multiplication: $I = V \cdot G$
 - Addition: $I^{column} = I_1^{row} + I_2^{row} + I_3^{row}$
- Ideal Nanoscale Devices for G :
 - Programmable Conductance
 - Multi-Level Cell
 - Small Footprint/High Density
 - Compatible with Existing CMOS Process

Memristors for In-Memory Computing

Also Called **Resistive Random Access Memory**, RRAM or ReRAM

Cross-section TEM



Programmable resistor w/ analog states

ISSCC: Intel adds embedded ReRAM to 22nm portfolio

January 03, 2019

TSMC to start embedded RRAM production in 2019

According to reports, Taiwan Semiconductor Manufacturing Company (TSMC) is aiming to start producing embedded RRAM chips in 2019 using a 22 nm process. This will be initial "risk production" to gauge market reception.

	Multi-Level Cell	Cell Area	R/W Speed
SRAM	×	large	Fast
DRAM	×	medium	Medium
1T1R	✓	medium	Medium Fast
Flash	✓	small	Slow




My work: Emerging Memory-Centric Design

- **Circuits & Systems Implementation**

- Spike-based Interface [DAC'15, DAC'18, DAC'20]
- Implementation of Neural Networks [VLSI'19, DAC'20]

- **Tolerate/Exploit Non-ideal Behavior of Memristors**

- Device Nonlinearity [ISCAS'16, IEDM'17, IEDM'19]
- Read Disturbance [ICCAD'17], Hard Fault [ITC'19]





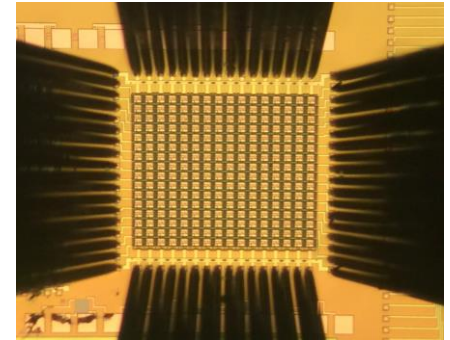
My work: Emerging Memory-Centric Design

- **Circuits & Systems Implementation**

- Spike-based Interface [DAC'15, DAC'18, DAC'20]
- Implementation of Neural Networks [VLSI'19, DAC'20]

- **Tolerate/Exploit Non-ideal Behavior of Memristors**

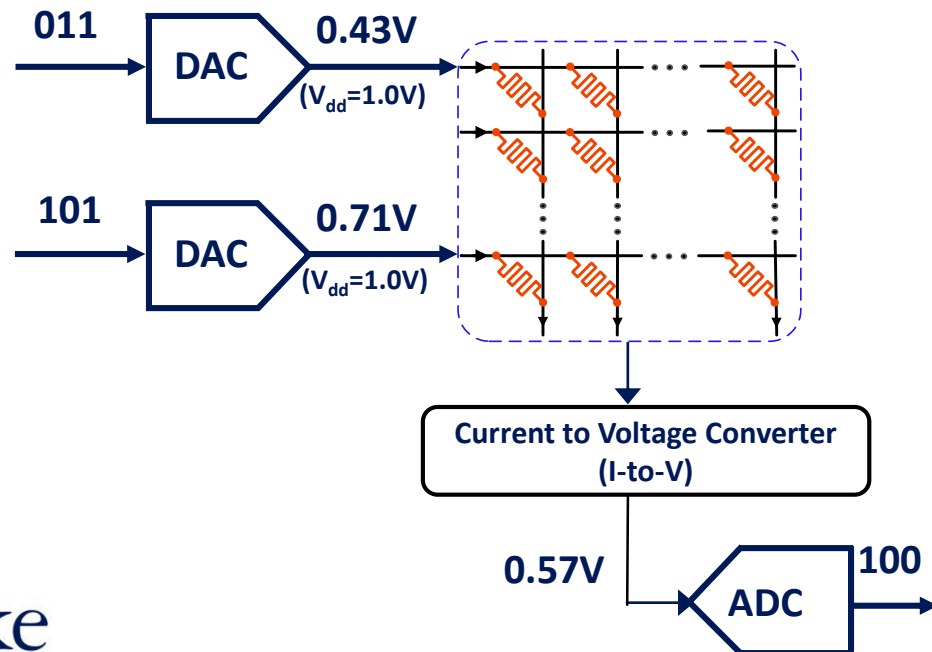
- Device Nonlinearity [ISCAS'16, IEDM'17, IEDM'19]
- Read Disturbance [ICCAD'17], Hard Fault [ITC'19]



Conventional ADC is Too Large

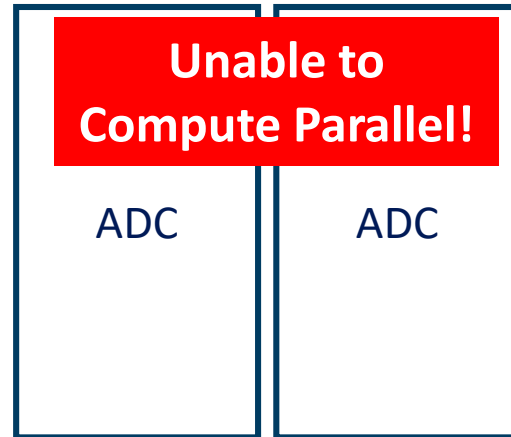
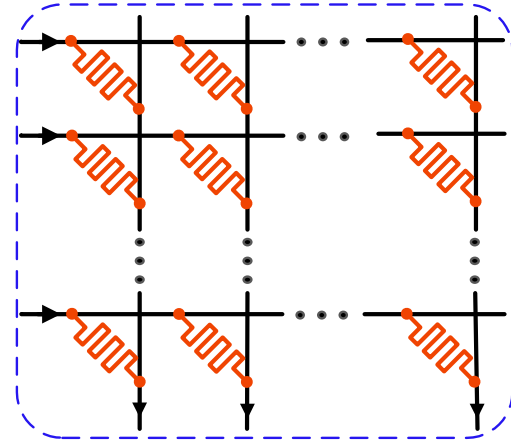
The Level-based Design

- Compatible to existing signal processing
- High speed computation



Actual Layout:

256x256 1T1R Array

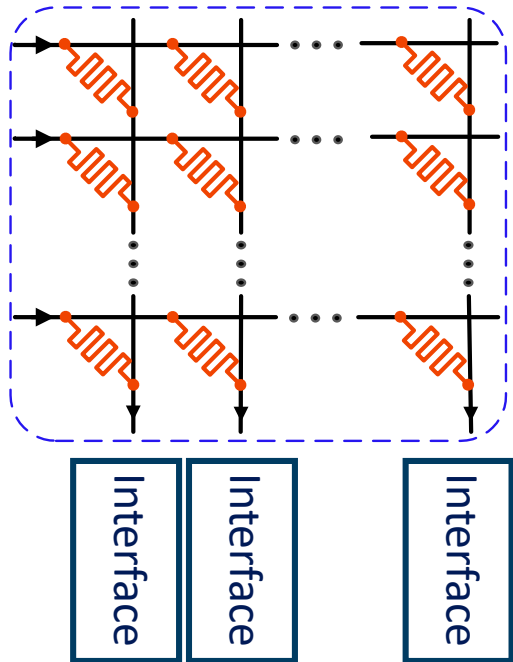


Based on 1.66MF² 8bit ADC
by K. Ohhata (JSSC 2019)

My Approach: Spiking Interface Circuit


What Better Designs Look Like

- Compute Parallely (Massive)
- Need Light-Weight Interface Circuitry


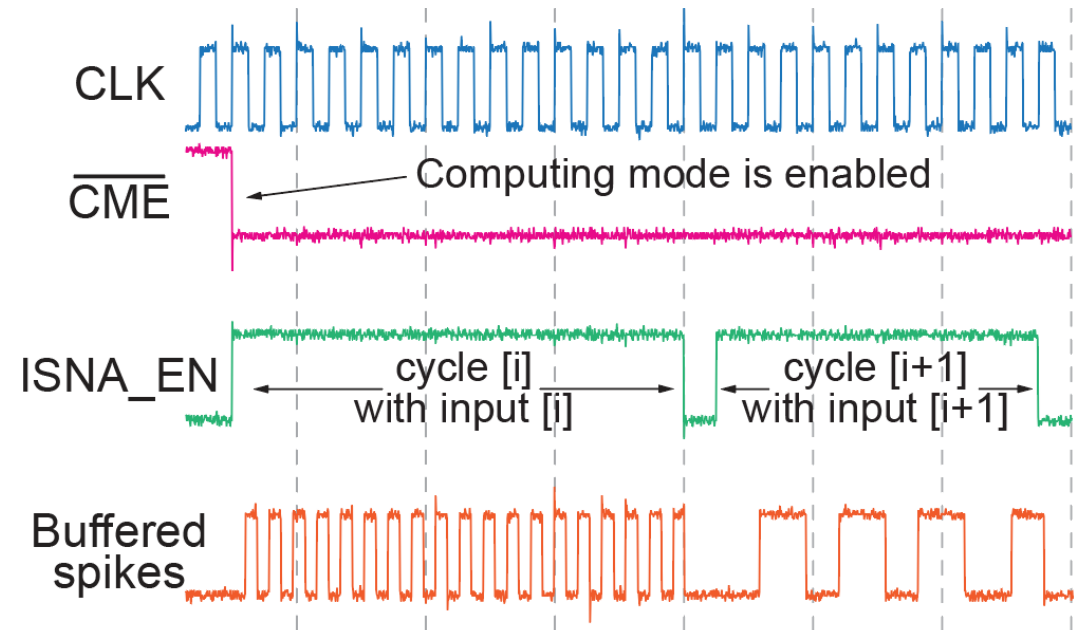
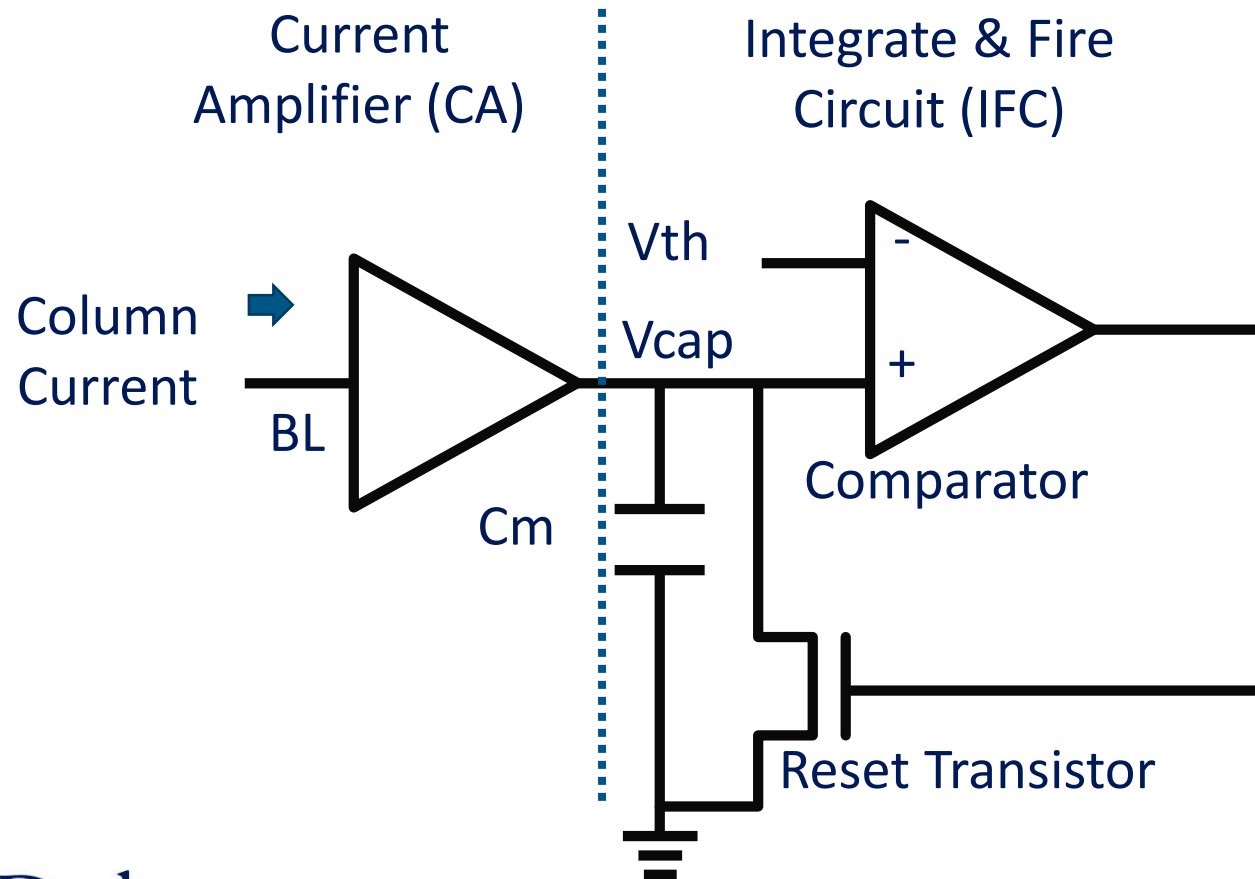


The Spike-based Design


- Closer to biological system
- Extremely high power efficiency




Spike Conversion




Spike Conversion
Circuit & Controller
3152x3152 μm^2




Spike Conversion



- 5.12 GOPS @ 8-bit precision
- 200ns latency @ 8-bit precision
- 43x area reduction vs. ADC by K. Ohhata (JSSC 2019)



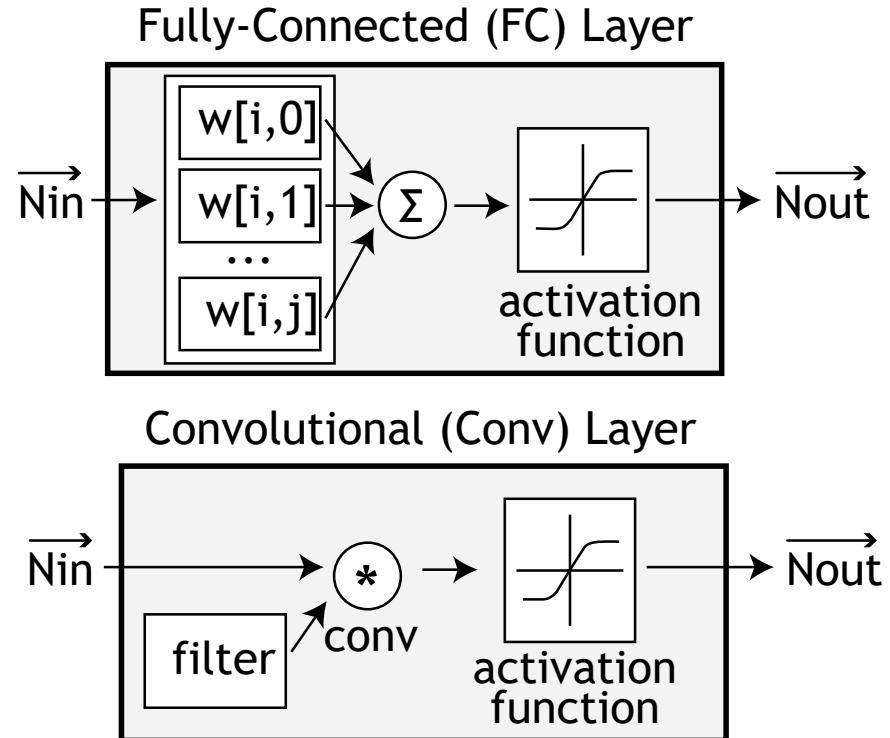
Tradeoff between large input current range and response speed



Enlarge phase margin tolerating capacitor positive feedback

How to Use Spiking-Based Design to Execute Neural Networks?

In Situ Nonlinear Activation (ISNA) Function




Single-layer Inference Operation:

- D** • Step 1: Load data from buffer
- A** • Step 2: Vector-matrix multiplication
- D** • Step 3: Nonlinear activation function
- D** • Step 4: Pooling
- D** • Step 5: Store results to buffer

D : digital domain **A** : analog domain

In Situ Nonlinear Activation (ISNA) Function



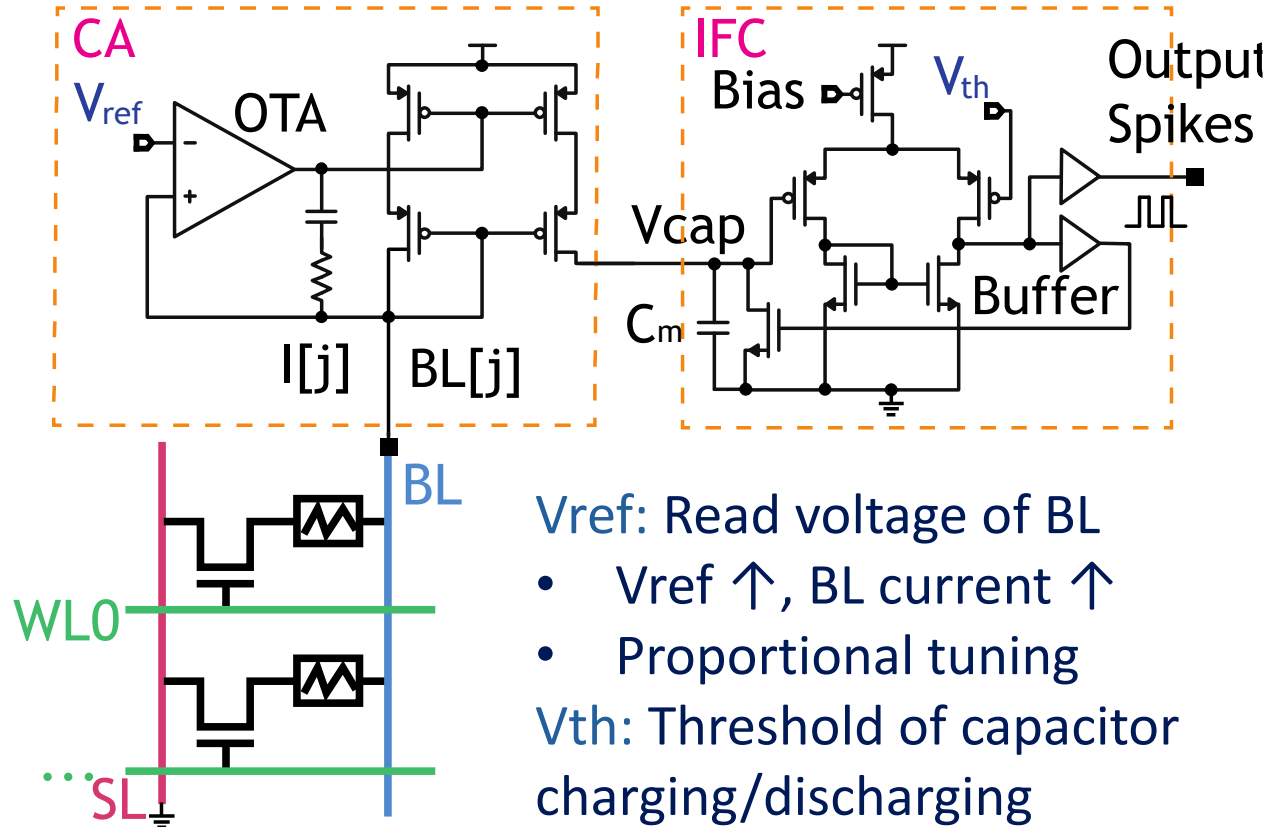
Single-layer Inference Operation:

- D** • Step 1: Load data from buffer
- A** • Step 2: Vector-matrix multiplication
- A** • Step 3: Nonlinear activation function
- D** • Step 4: Pooling
- D** • Step 5: Store results to buffer

D : digital domain **A** : analog domain

Combine Step 2 & Step 3 to simplify PE operation:
Use linear + nonlinear regions

Adjust Activation Function




V_{ref} : Read voltage of BL


- $V_{ref} \uparrow$, BL current \uparrow
- Proportional tuning

V_{th} : Threshold of capacitor charging/discharging

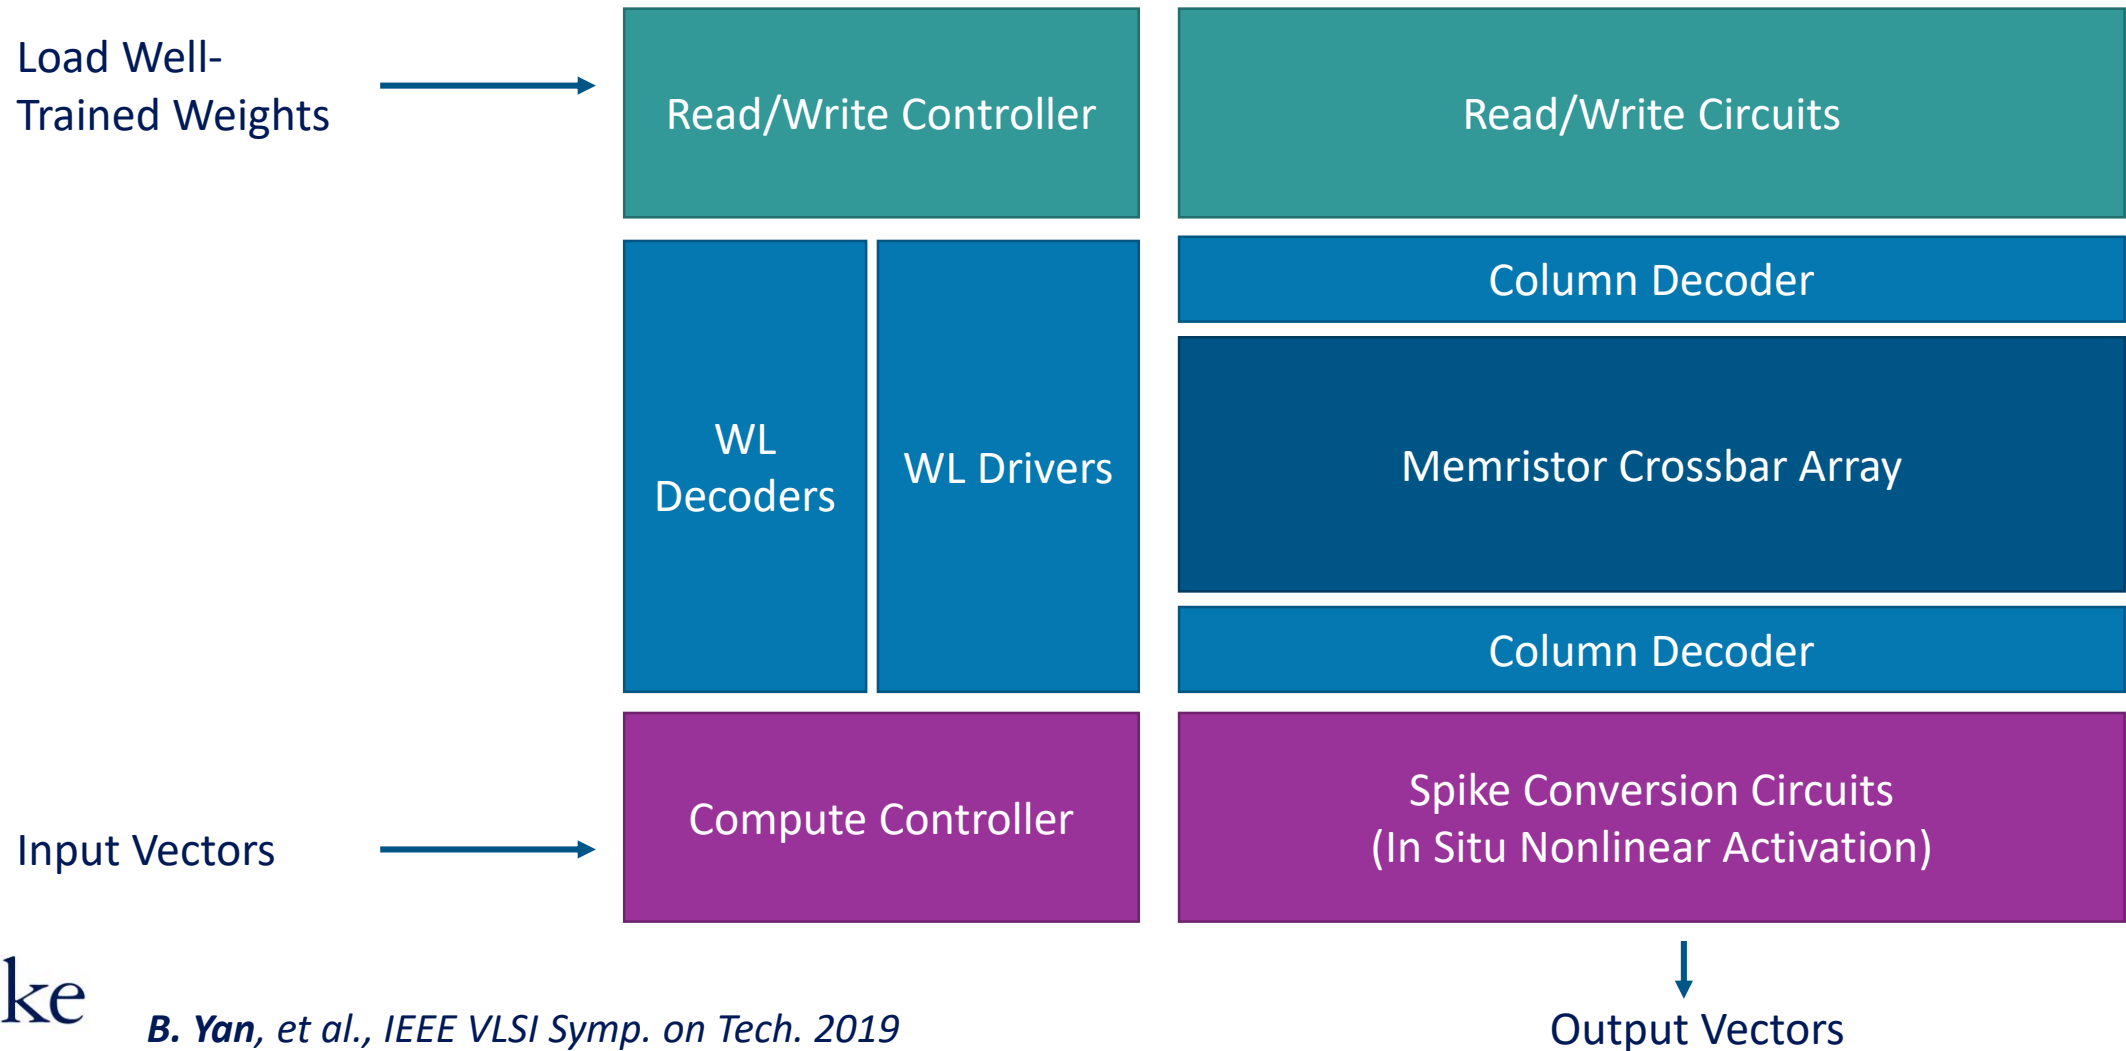
- $V_{th} \downarrow$, Charging/discharging \uparrow
- Distorted tuning




Function clipped-relu



Chip Architecture




Chip Summary



Technology	150nm CMOS +HfO _x RRAM
Macro Capacity	64K (256×256)
Clock Frequency	50MHz
Energy Efficiency	0.257pJ/Mac
Average Power	1.52 mW
Layer-wise Latency	200ns
Real-time Benchmarks	3-layer perceptrons, LeNet-4, LetNet-5



GUI



Die Photo of RRAM Neuro-computing Engine



Camera

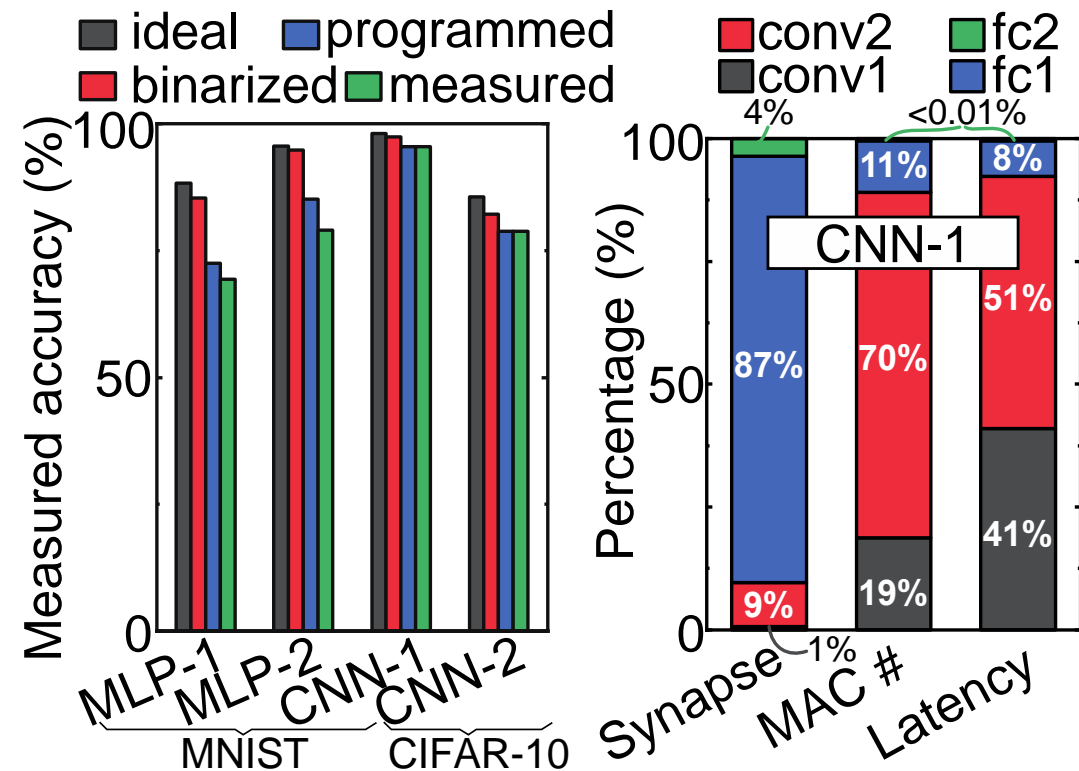
Evaluation: Measured Neural Network Results

MNIST:

- MLP-1: Single-layer perceptron
- MLP-2: 2-layer perceptron
- CNN-1: 4-Layer LeNet




CIFAR-10:

- CNN-2: 5-Layer LeNet





My work: Emerging Memory-Centric Design

- Circuits & Systems Implementation
 - Spike-based Interface [DAC'15, DAC'18, DAC'20]
 - Implementation of Neural Networks [VLSI'19, DAC'20]
- **Tolerate/Exploit Non-ideal Behavior of Memristors**
 - Device Nonlinearity [ISCAS'16, IEDM'17, IEDM'19]
 - Read Disturbance [ICCAD'17], Hard Faults [ITC'19]



Non-Ideal Memristor - I



- Cell Nonlinearity: conductance varies when applied with different voltages



- Solution: Current Amplifier to Clamp Cell Voltage (shown in previous circuit design part)

Non-Ideal Memristor - II


- Memristance Drift: conductance/memristance gradually deviates from original values under read voltage (read disturbance)
- Characteristic:
 - Very slow to observe



How to Mitigate Memristance Drift for Inference?

$$E_{mse} = \sum_j^n \left(t_j - \sum_i^m w_{ij} x_i \right)^2 \quad E'_{mse} = \sum_j^n \left(t_j - \sum_i^m w_{ij} x_i - \sum_i^m \Delta w_{ij} x_i \right)^2$$

E_{MSE} : MSE Error Function
 t_j : label for each neuron

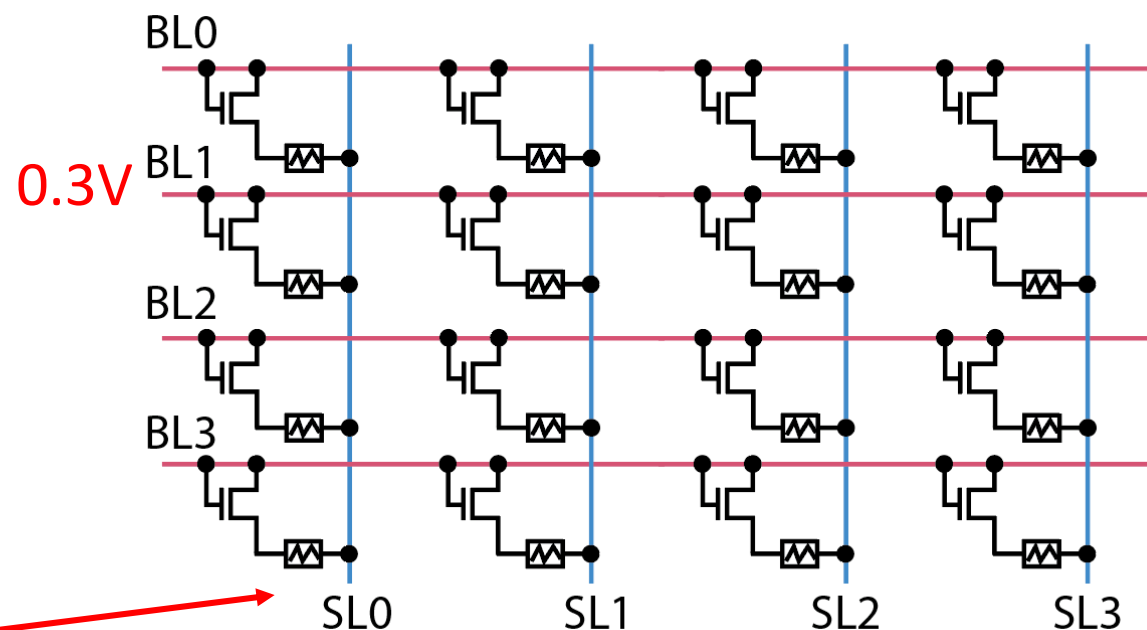


$$\begin{aligned} \Delta E_{mse} &= E'_{mse} - E_{mse} \\ &= -2 \sum_j^n \left[\left(t_j - \sum_i^m w_{ij} x_i \right) \left(\sum_i^m \Delta w_{ij} x_i \right) \right] \end{aligned}$$

Minimize Degraded Error Function:

$$\frac{\partial \Delta E_{mse}}{\partial \Delta w_{ij}} = -2 \left(t_j - \sum_i^m w_{ij} x_i \right) x_i \leq 0$$

Duke

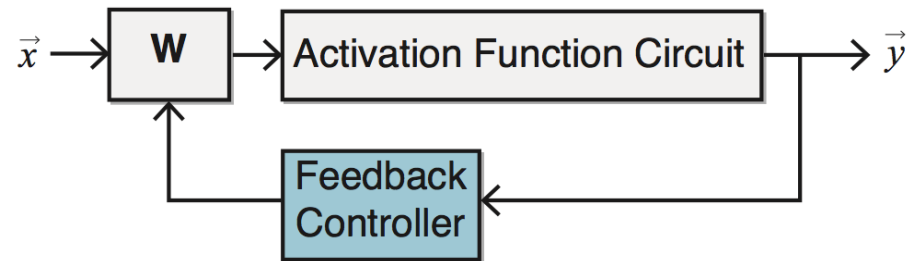


Change Column Current Direction



0V/0.6V

Closed-loop Design to Enhance Weight Stability

Feedback controller: Adjust the voltage condition to compensate the memristance drift.



“Arrogant principle”: last output is used as the label to determine the sense condition.

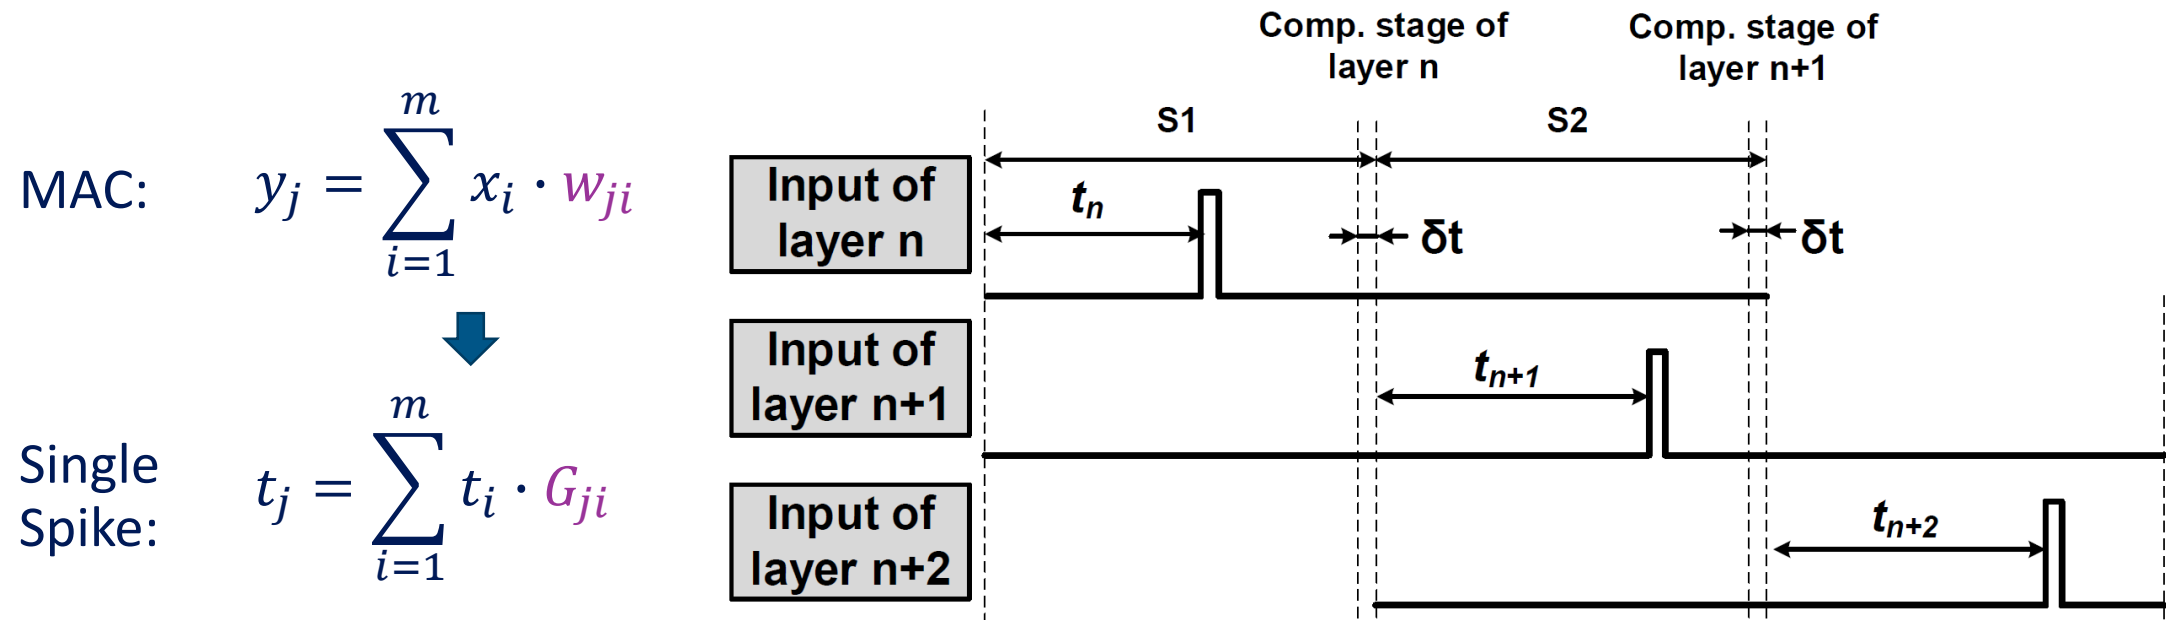


Summary

- The Past and Now of AI Hardware
 - In-Memory Computing Eliminates Weights Movement
- My Work:
 - Spiking-based In-Memory Computing Engine Offer Very High Energy Efficiency & Good Performance
 - Clever Design Methodologies (e.g., Closed-Loop Sensing) is Effective to Tolerate Nonideal Features of Memristors

Future Work I: Single Spike Processing Engine

Use Single Spike to Replace Multiple Spikes: ~10x Improvement of Energy Efficiency



“ReSiPE: ReRAM-based Single-Spiking Processing-In-Memory Engine”

Simulation Results Coming in July at Design Automation Conference (DAC) 2020


Duke

Expect to IC Tape-Out + Single-Spike Encoding Method [Potential Collaboration Opportunity]

Future Work II: In-Memory Compute & Cache

Challenge:

During Training:

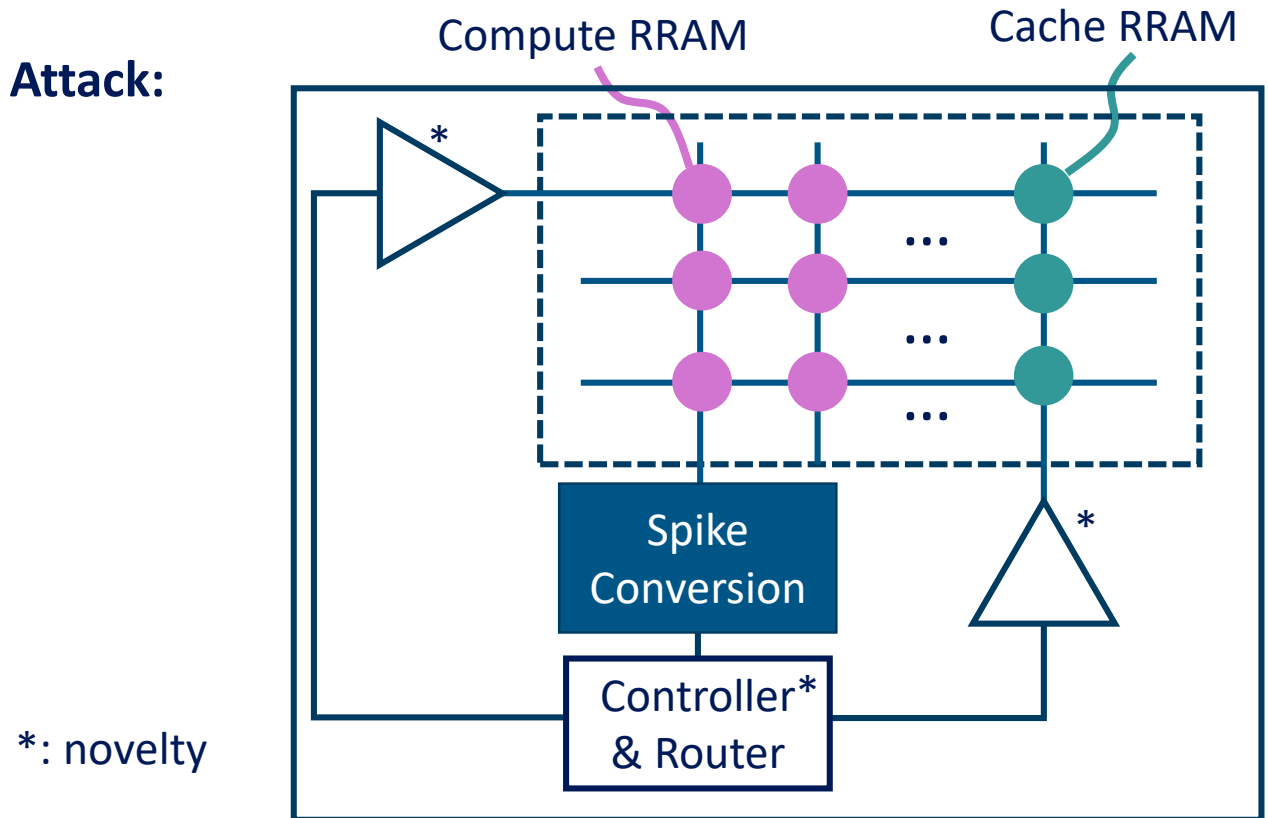


No Need to
Transport Weight

External DDRAM
Access (Expensive!)

Duke

Attack:



*: novelty

Prospect Applications:


NN Training Acceleration

Self-Updatable In-Memory Computing Engine


Long-Term Prospects

Compute


Rethink Compute Memory Hierarchies



RRAM




PCM



MRAM


“Organ” for Compute &
Cache w/ different
emerging memory types

In Situ Data Processing (Near-Sensor Computing)



Control

Analog Content Addressable Memories



Fuzzy Logic Hardware

Many Thanks for the Support!

Advisors:

Prof. **Hai “Helen” Li**

Prof. **Yiran Chen**



Collaborators:

Prof. **Joe Qiu**

Army Research Office

Dr. **Qing Wu**

Air Force Research Laboratory

Prof. **Jianhua (Joshua) Yang**

UMASS Amherst

Prof. **Meng-Fan Chang**
National Tsing-Hua Univ.



Prof. **Krishnendu Chakrabarty**

Duke University

Prof. **Weisheng Zhao**
Beihang University

Student Collaborators:

Ziru Li, Qilin Zheng, Brady Taylor

Duke



Qualcomm



Thanks for Listening & Qs!

slides available at:

<https://bonanyan.github.io/bn/>